

# MAKERERE



# UNIVERSITY

P.O. Box 7062 Kampala Uganda

<https://www.cs.mak.ac.ug>

Telephone: 256-414-534560/1-9

E-mail: [cs@cis.mak.ac.ug](mailto:cs@cis.mak.ac.ug)

## **COLLEGE OF COMPUTING AND INFORMATION SCIENCES**

### Smart Traffic Management: An AI-Driven Solution for Optimizing Traffic Flow in Developing Countries

by

Tusiime Fortunate

Ssali Joshua

Department of Computer Science  
School of Computing & Informatics Technology

#### **Supervisors**

Dr. Mary Nsabagwa

Dr. Brian Muchake

May 2025

# Contents

<b>1</b>	<b>General Introduction</b>	<b>1</b>
1.1	Introduction . . . . .	1
1.2	Background . . . . .	2
1.3	Problem Statement . . . . .	3
1.4	Objectives . . . . .	4
1.4.1	Main Objective . . . . .	4
1.4.2	Specific Objectives . . . . .	4
1.5	Scope . . . . .	4
1.5.1	Geographical Scope . . . . .	5
1.5.2	Content Scope . . . . .	5
1.5.3	Time Scope . . . . .	6
1.6	Significance of the Study . . . . .	6
<b>2</b>	<b>Literature Review</b>	<b>8</b>
2.1	Traffic Management Systems . . . . .	8
2.2	AI-Driven Solutions . . . . .	8
2.3	Hardware-Based Solutions . . . . .	9
2.4	You Only Look Once (YOLO) . . . . .	10
2.5	Manual Solutions . . . . .	11
2.6	Comparison with Our Solution . . . . .	12
<b>3</b>	<b>Methodology</b>	<b>14</b>
3.1	Research Design . . . . .	14
3.2	Data Gathering . . . . .	14
3.3	Data Preparation . . . . .	14
3.4	Model Design . . . . .	15
3.5	Simulation . . . . .	16
3.6	Overall System Architecture . . . . .	16
3.6.1	Master Slave Architecture . . . . .	16

3.6.2	Slave Microcontroller . . . . .	17
3.6.3	Master Microcontroller . . . . .	18
3.6.4	System Architecture . . . . .	20
<b>4</b>	<b>Data Collection and Preparation</b>	<b>22</b>
4.1	Data Gathering . . . . .	22
4.2	Data Collecting . . . . .	22
4.3	Data Preparation . . . . .	23
4.3.1	Preprocessing for Traffic Classification . . . . .	23
4.3.2	Preprocessing for Vehicle Detection . . . . .	24
4.3.3	Dataset Splitting . . . . .	24
4.3.4	Quality Assurance . . . . .	25
4.4	Results of Data Collected . . . . .	25
<b>5</b>	<b>Model Design and Training</b>	<b>27</b>
5.1	Model Selection . . . . .	27
5.2	Model Architectures . . . . .	27
5.2.1	EfficientNetB0 . . . . .	27
5.2.2	YOLOv8 . . . . .	29
5.3	Training Process . . . . .	29
<b>6</b>	<b>Evaluation</b>	<b>31</b>
6.1	Traffic Classification Model . . . . .	31
6.2	Car Detection Model . . . . .	32
6.3	Simulation Validation . . . . .	34
<b>7</b>	<b>Discussion of Results</b>	<b>36</b>
7.1	Traffic Classification Performance . . . . .	36
7.2	Vehicle Detection Accuracy . . . . .	37
7.3	Simulation Validation . . . . .	38
7.4	Comparison to Existing Approaches . . . . .	39

<b>8</b>	<b>Future Work</b>	<b>40</b>
8.1	Hardware and Deployment . . . . .	40
8.1.1	Microcontroller Selection . . . . .	40
8.1.2	Model Optimization for Hardware . . . . .	40
8.1.3	Integration with Traffic Infrastructure . . . . .	41
8.1.4	Testing in Controlled Environments . . . . .	41
8.1.5	Pilot Deployment . . . . .	41
8.2	System Validation and Scalability . . . . .	41
8.2.1	Pilot Studies . . . . .	41
8.2.2	Scalable Architecture . . . . .	42
8.2.3	Data-Driven Improvements . . . . .	42
8.2.4	Continuous Model Maintenance . . . . .	42
8.3	Technical Enhancements . . . . .	42
8.3.1	Multi-Modal Data Integration . . . . .	42
8.3.2	Adaptive Learning . . . . .	43
8.3.3	Non-Vehicle Traffic Integration . . . . .	43
8.3.4	Predictive Analytics . . . . .	43
8.3.5	Advanced Model Exploration . . . . .	43
8.3.6	Anomaly Detection . . . . .	43
8.3.7	Handling Adverse Conditions . . . . .	44
8.4	Research and Development . . . . .	44
8.4.1	Explainability and Trust . . . . .	44
8.4.2	Smart City Integration . . . . .	44
8.4.3	Standardization and Collaboration . . . . .	44
8.5	Ethical and Societal Considerations . . . . .	44
8.5.1	Equity in Traffic Management . . . . .	45
8.5.2	Safety and Reliability . . . . .	45
8.5.3	Data Privacy and Security . . . . .	45
8.5.4	Environmental Impact . . . . .	45

<b>9 Conclusion</b>	<b>46</b>
9.1 Research Achievements . . . . .	46
9.2 Implications for Urban Traffic Management . . . . .	47
9.3 Comparison to Existing Approaches . . . . .	48
9.4 Limitations and Future Directions . . . . .	48
9.5 Significance and Broader Impact . . . . .	49
9.6 Final Remarks . . . . .	50
<b>Appendices</b>	<b>58</b>
<b>Appendix A Data Collection</b>	<b>58</b>
A.1 Kaggle Datasets . . . . .	58
A.2 Techniques for Gathering Opinions . . . . .	58
<b>Appendix B Causes of Traffic Problems in Developing Countries</b>	<b>59</b>
<b>Appendix C Methodology Justification</b>	<b>61</b>
<b>Appendix D GitHub Repository Details</b>	<b>61</b>
D.1 Repository Overview . . . . .	61
D.2 Repository Structure . . . . .	61
D.3 Setup Instructions . . . . .	62

# 1 General Introduction

This paper presents a traffic management system designed to classify and detect traffic levels in real-time using machine learning and computer vision techniques, tested and validated through simulations in SUMO (Simulation of Urban Mobility) [1]. The system employed the EfficientNetB0 model for classifying traffic into high, low, and medium categories, achieving superior performance with a macro-averaged F1-score of 0.91 and accuracy of 0.91. For vehicle detection and counting, the YOLOv8x model was used because of its high inference speed up to 100 Frames Per Second and accuracy [2], as demonstrated by consistent car counts with confidence scores ranging from 0.49 to 0.95 across diverse traffic scenes. The solution was tested in SUMO, simulating dynamic urban traffic scenarios, where it effectively optimized traffic flow by adapting to varying densities and conditions. The system was optimized for deployment on a hardware microcontrollers which use the serial peripheral communication protocol to share this data, analyze it using the models and then manage traffic. This approach highlights the ability to manage traffic in a real world setting while offering a scalable and efficient solution to urban congestion challenges.

## 1.1 Introduction

Traffic congestion in developing countries is a significant challenge which causes problems such as increased travel times, fuel wastage, environmental degradation, and safety risks such as theft during heavy traffic. In this paper, we propose an innovative smart traffic management system that leverages Machine Learning and data science to optimize traffic flow. Our system dynamically allocated green light time to lanes based on real-time vehicle counts, aiming to reduce congestion, enhance fuel efficiency, and improve road safety. This project sought to develop and test a cost effective solution that suited the budget constraints of developing countries, ensuring scalability and sustainability. By implementing this solution, we aim to contribute to a more efficient and environmentally friendly urban transport sector.

## 1.2 Background

Rapid urbanization across developing countries has led to a significant increase in the number of vehicles on the road, intensifying the problem of traffic congestion. In many urban areas, commuters spend hours in traffic daily, resulting in economic losses due to wasted fuel, reduced productivity, and increased pollution levels [3]. Smart traffic management refers to the application of advanced technologies—especially Artificial Intelligence (AI)—to monitor, analyze, and manage traffic conditions in real time. One of its key objectives is to achieve optimized traffic flow, which entails the efficient movement of vehicles through road networks with minimal delays. Optimized traffic flow is accomplished by dynamically adjusting traffic signals, detecting congestion early, and adapting to changing road conditions using real-time data inputs. Conventional traffic control systems rely on fixed or pre-programmed signal cycles. These static systems fail to respond to real-time fluctuations in traffic volume, often leading to inefficient road use and long vehicle queues. More recent solutions have employed induction loop detectors, which are embedded beneath road surfaces to detect vehicles by sensing disturbances in magnetic fields caused by metal objects [4]. While these systems can improve responsiveness, their installation and maintenance are expensive and technically demanding, posing challenges for implementation in resource-constrained environments. Earlier efforts in traffic control have also included manual oversight and time-based loop systems. However, these methods are inadequate for addressing the complexity and variability of modern traffic patterns, especially during peak hours or unexpected surges [5]. This paper proposes an AI-driven computer vision system as a smart, adaptable, and cost-effective traffic management solution. By leveraging real-time video data, the system can detect vehicles, analyze traffic density, and dynamically adjust traffic signals. The benefits of such an approach include reduced congestion, lower fuel consumption, decreased carbon emissions, improved road safety, and enhanced commuter experience—critical improvements for rapidly growing urban centers in developing regions.

### 1.3 Problem Statement

Traffic congestion in urban areas of developing countries significantly hinders economic development, environmental sustainability, and public safety due to rapid urbanization, rising vehicle ownership, and inadequate transportation infrastructure. With projections indicating that over 50% of Africa’s population will live in urban areas by 2035, the surge in vehicle numbers, particularly during peak hours, overwhelms narrow, poorly maintained road networks, made worse by a heavy reliance on private vehicles due to underdeveloped or unreliable public transportation systems [6, 7]. These urban areas also use static traffic management systems like fixed-cycle traffic lights, which fail to adapt to dynamic traffic patterns, causing significant delays at busy intersections, while poor urban planning such as inadequate road networks and a lack of integrated transport policies, further aggravates congestion [8, 9]. Informal transport modes like motorcycles and unregulated minibus taxis contribute to chaotic traffic flow due to their unpredictable movements, and limited investment in traffic management infrastructure, together with a shortage of trained personnel, hampers effective mitigation efforts. [10, 11]. Economically, congestion leads to prolonged travel times, reducing productivity and increasing transportation costs, with losses estimated at 2-5% of GDP annually in urban areas [8]. Environmentally, idling vehicles increase fuel consumption, contributing to greenhouse gas emissions and air pollution that often exceed WHO safety thresholds, resulting in health issues like respiratory diseases [12]. Socially, congestion heightens commuter stress, diminishes quality of life, and exacerbates inequality, as low-income groups face longer commutes due to limited access to efficient transport options [7]. Safety is also compromised, with erratic stop-and-go driving and chaotic traffic patterns linked to higher accident rates, particularly affecting vulnerable road users like pedestrians and motorcyclists [13]. Current solutions, such as manual traffic control by officers or volunteers, are inefficient and contribute to disorder, while advanced smart traffic systems used in developed nations are often too expensive and technically complex for resource-constrained developing countries [8]. This paper proposes a cost-effective, scalable, and adaptive traffic management system to address these multifaceted challenges and improve urban mobility in developing countries.



## 1.4 Objectives

### 1.4.1 Main Objective

The primary objective of this research was to develop a real time traffic management system that integrates machine learning techniques for traffic level classification and vehicle detection. The system was designed to facilitate automated decision-making for traffic control, such as optimizing traffic signal timings or rerouting vehicles, and was optimized for deployment on a resource constrained hardware microcontroller hence reducing urban traffic congestion, enhancing safety, and improving transportation efficiency in modern cities.

### 1.4.2 Specific Objectives

1. To gather images from city traffic cameras, capturing various traffic conditions, and obtain additional data from more developing countries hence ensuring a diverse dataset for training the AI models.
2. To preprocess and prepare the collected data to ensure it is compatible with machine learning models.
3. To train multiple machine learning models for traffic classification, select the best performer , and implement the vehicle detection model as either YOLOv8x or RCNN(Region Based Convolutionary Networks ).
4. To test and validate the algorithms using SUMO (Simulation of Urban Mobility) simulations to measure their ability to optimize traffic flow under various urban scenarios such as peak hours and congested conditions.

## 1.5 Scope

The scope of the research is defined across three dimensions to provide clear boundaries and focus for the study:

### 1.5.1 Geographical Scope

This study focuses on Uganda, with a particular emphasis on Kampala, the capital city and primary economic hub, where traffic congestion poses significant challenges to urban mobility and sustainable development. Uganda, located in East Africa between longitudes 29E to 35.2E and latitudes 1.5S to 4.5N, spans approximately 241,038 km<sup>2</sup> and is bordered by Kenya, South Sudan, the Democratic Republic of Congo, Rwanda, and Tanzania [14]. Kampala, situated 8 km north of Lake Victoria and covering 190 km<sup>2</sup>, is Uganda's administrative, commercial, and cultural center, divided into five divisions: Central, Kawempe, Makindye, Nakawa, and Lubaga [15]. The city's population, estimated at 1.7 million in 2020, is growing at a rate of 4.1% annually, making it one of Africa's fastest-growing urban areas [6]. This rapid urbanization, coupled with a road network that has not expanded commensurately, results in severe traffic congestion, particularly during peak hours [16]. Kampala's transportation system is dominated by private vehicles, informal public transport modes like boda-bodas (motorcycles) and minibus taxis, and heavy-duty vehicles, exacerbated by inadequate public transport infrastructure [17]. The city's strategic location along the Northern Corridor, a key transport route connecting Mombasa to Rwanda and beyond, amplifies the regional impact of its traffic bottlenecks [18]. Recent initiatives by the Kampala Capital City Authority (KCCA), supported by the Japanese International Cooperation Agency (JICA), include the development of a Traffic Control Center (TCC) launched in 2022, equipped with real-time monitoring and smart signal systems to mitigate congestion [19]. However, challenges such as poor driver behavior, limited traffic enforcement, and insufficient road capacity persist, necessitating innovative solutions like the cost-effective, machine learning-based traffic management system proposed in this study [20].

### 1.5.2 Content Scope

This research encompasses traffic classification using both deep learning and traditional machine learning models, real-time vehicle detection with YOLOv8x, and the development of decision-making algorithms for automated traffic management. It also includes the integration of all components into a microcontroller-compatible system, validation through SUMO-

based simulations, and model optimization for embedded deployment. The study is limited to simulated environments and urban traffic scenarios, excluding real-world deployment and rural or highway traffic contexts.

### 1.5.3 Time Scope

The research is conducted over a one-year period, from May 2024 to May 2025. This time-frame encompasses all phases of the project, as shown in the table below:

Table 1.1: A table showing the expected time to be taken achieving the different phases of the project

Phase	Duration	Description
Data Collection and cleaning	May 2024 to June 2024	Collecting data, interviewing stakeholders, analyzing and cleaning the data.
Model Development and Evaluation	July 2024 to August 2024	Training, testing, and selecting the optimal models for traffic classification and vehicle detection.
System Implementation	September 2024 to November 2024	Developing decision-making algorithms and integrating components.
Simulation and Validation	December 2024 to February 2025	Conducting SUMO simulations to assess system performance.
Finalization and documentation	March 2026 to May 2026	Optimizing models for microcontroller deployment and documenting the work done.

## 1.6 Significance of the Study

For **road users**, the system significantly enhances daily commuting by optimizing traffic flow, reducing travel times, and minimizing delays caused by inefficient static traffic lights or manual interventions, which currently lead to congestion especially during peak hours.

This leads to improved quality of life, particularly for low income commuters who face disproportionate commuting burdens, and enhances public safety by reducing accident risks associated with chaotic traffic patterns. For the **government** bodies, this system offers a practical and budget-friendly alternative to expensive smart traffic systems used in developed nations, enabling data driven urban planning and integration with existing initiatives like the KCCA-JICA Traffic Control Center. Economically, it tackles congestion-related losses—estimated at 2–5% of GDP in urban areas of developing countries—boosting productivity by streamlining traffic movement. Environmentally, it supports sustainability goals like SDG 11 (Sustainable Cities and Communities) by reducing fuel consumption and greenhouse gas emissions. For **enforcement officers**, the system provides real-time, accurate traffic data through AI-driven detection, enabling more effective traffic control, reducing the reliance on manual interventions, and improving coordination to maintain order on the roads . Academically, this study advances the application of machine learning in resource-constrained environments, offering insights into AI model optimization and fostering scalable solutions for other developing countries.

## 2 Literature Review

### 2.1 Traffic Management Systems

Traffic management systems (TMS) are designed to enhance the efficiency and safety of transportation networks by integrating communication, sensing, and processing technologies. A review [21] classifies TMS into infrastructure-based and infrastructure-free systems, detailing their architecture and operational phases. TMS typically involve three stages: information gathering from vehicles (e.g., GPS, speedometers), sensors (e.g., traffic lights, road occupancy), and web sources; information processing to identify hazards like congestion or accidents; and service delivery, such as route suggestions or speed adjustments. Examples include cooperative congestion detection systems like CoTEC and infrastructure-based solutions like EcoTrec, which optimize traffic flow through centralized traffic management centers (TMCs). These systems leverage vehicle-to-vehicle (V2V) and vehicle-to-infrastructure (V2I) communications, using protocols like Dedicated Short-Range Communication (DSRC) and LTE. Challenges include heterogeneous data integration, big data management, and security concerns, particularly in vehicular ad-hoc networks (VANETs). Another study by Mittal et al. (2020) explores neuro-fuzzy-based traffic light management, demonstrating improved traffic flow through adaptive signal timing, though scalability remains a concern in dense urban settings.

A review [22] provides an update of intelligent traffic management systems, highlighting the integration of AI and IoT technologies to enhance real-time decision-making and traffic efficiency. Additionally, the Federal Highway Administration [23] emphasizes the need for modern TMS to adapt to increasing traffic volumes and leverage real-time data processing for improved performance.

### 2.2 AI-Driven Solutions

AI and machine learning have transformed traffic management by enabling predictive analytics and real-time decision-making. A survey by Manibardo et al. (2021) examines deep learning applications in congestion detection, prediction, and alleviation, highlighting mod-

els like convolutional neural networks (CNNs) and recurrent neural networks (RNNs). These models analyze large datasets from traffic cameras and sensors to forecast congestion and optimize traffic signal timings. For example, reinforcement learning algorithms dynamically adjust traffic lights based on real-time traffic flow, reducing delays by up to 20% in simulated urban scenarios. Another study by Akhtar and Moridpour (2021) reviews AI-based traffic congestion prediction, noting that short-term forecasting using historical and real-time data achieves high accuracy with models like Long Short-Term Memory (LSTM) networks. A practical implementation in Pécs, Hungary, by Hajgató et al. (2023), used LSTM for public bus speed prediction, detecting congestion with 82.37% accuracy using incremental learning, which is more scalable than traditional batch learning. However, challenges include the need for large, high-quality datasets and computational complexity, which can hinder real-time applications in resource-constrained environments.

A recent review [24] discusses advances in AI-based traffic prediction models, focusing on multivariate traffic time series modeling and the use of federated learning for privacy-preserving data analysis.

## 2.3 Hardware-Based Solutions

Hardware components are critical for collecting and processing traffic data in intelligent traffic management systems (ITMS). A review by Alam et al. (2023) details hardware used in ITMS, including image sensors (e.g., Sony’s 127.68-megapixel sensors), CMOS sensors (e.g., Python XK at 80 FPS), GPS devices, RFID-based toll systems, and lidar/radar sensors. These technologies enable real-time traffic monitoring, vehicle detection, and incident management. For instance, Hikvision and Citilog provide advanced video surveillance systems for traffic flow analysis, while GPS sensors integrated into applications like Waze and Google Maps offer real-time navigation updates. Edge computing devices, as discussed by Intellias (2022), process data locally to reduce latency, enhancing system responsiveness. However, challenges include high installation and maintenance costs, as well as the need for robust coordination among multiple cameras and sensors to handle varying environmental conditions like weather and illumination changes.

Recent advancements in sensor technology [25], include the development of low-cost, high-

accuracy sensors that can operate in diverse environmental conditions, making them suitable for deployment in developing countries.

## 2.4 You Only Look Once (YOLO)

The YOLO model has emerged as a leading approach for real-time object detection. YOLO revolutionized object detection by framing it as a single regression problem, achieving high speed and accuracy compared to two-stage models like the Faster R-CNN [26]. YOLO's single-stage architecture processes images in one pass, predicting bounding boxes and class probabilities directly, which enables real-time performance critical for traffic applications. The YOLOv3 was evaluated against Faster R-CNN and Single Shot MultiBox Detector (SSD), this showed that the YOLOv3 achieved a mean Average Precision (mAP) of 57.9% on the COCO dataset at 20 frames per second (FPS), outperforming SSD (mAP: 41.2%) and approaching Faster R-CNN's accuracy (mAP: 59.1%) at significantly higher speeds [27]. The evolution of YOLO has led to significant improvements. YOLOv5 introduced optimizations like mosaic data augmentation and auto-learning bounding box anchors, achieving an mAP of 68.2% on COCO at 140 FPS on a V100 GPU [28]. YOLOv8, released by Ultralytics in 2023, further enhancing performance with an anchor free architecture and improved backbone (C2f module), achieving an mAP of 73.5% on COCO and 50 FPS on edge devices like Jetson Nano [29]. YOLOv8 was applied to urban traffic surveillance, reporting a 92.3% accuracy in vehicle detection under varying lighting and weather conditions, surpassing EfficientDet (mAP: 70.1%) and SSD (mAP: 65.4%) in real-time scenarios [30]. Compared to other object detection alternatives, YOLO's lightweight architecture and real-time capabilities make it ideal for resource constrained environments like in developing countries. Faster R-CNN, while accurate, is computationally intensive (7 FPS on high-end GPUs), limiting its edge deployment [31]. SSD offers speed but sacrifices accuracy in complex scenes [32]. EfficientDet balances accuracy and efficiency but requires more memory than YOLOv8x [33]. These attributes justify our selection of YOLOv8x for vehicle detection, as it aligns with the need for high accuracy, real-time performance, and cost-effective deployment in developing countries.

## 2.5 Manual Solutions

Manual solutions to traffic congestion rely on traditional methods such as road redesign, traffic police deployment, and policy interventions. According to Wikipedia contributors (2025), strategies like increasing road capacity through lane additions or junction improvements aim to enhance traffic flow, but they often induce demand, leading to increased traffic volumes, as noted in a 2011 study in *The American Economic Review*. Congestion pricing, implemented in cities like London and Singapore, charges drivers for entering congested zones, reducing traffic volumes by 16% in London (2006) and maintaining speeds of 20-30 km/h on Singapore’s arterial roads. Traffic police manage flow at intersections, particularly during peak hours or incidents, but their effectiveness is limited by human fatigue and scalability. Policy measures like parking restrictions and park-and-ride systems encourage public transport use, though their impact varies by city. A study by Anas and Lindsey (2023) critiques road capacity expansions, suggesting that urban planning and demand-side policies are more sustainable long-term solutions.

The study [34] evaluates the long-term effects of congestion pricing in major cities, finding that while initial reductions in traffic volumes are significant, sustained behavior change requires complementary policies such as improved public transportation.



## 2.6 Comparison with Our Solution

Table 2.1: Comparison of Traffic Management Solutions. Our proposed system integrates EfficientNetB0 for traffic classification and YOLOv8x for vehicle detection, optimized for edge microcontrollers.

Feature	TMS	AI-Based	Hardware-Based	Manual	Our System
Real-Time Decisions	✓	✓	✓		✓
Adaptive Signal Control	✓	✓		✓	✓
Edge Computing Support			✓		✓
High Accuracy Detection		✓	✓		✓
Cost-Effective Deployment				✓	✓
Scalability in Developing Countries					✓
Reduced Infrastructure Needs					✓
Low Computational Complexity				✓	✓
Uses Deep Learning Models		✓			✓

Our proposed system integrates EfficientNetB0 for traffic classification and YOLOv8x for vehicle detection, optimized for edge microcontrollers. Validated through SUMO simulations, it addresses imbalanced traffic data and complex patterns. Compared to TMS, our AI-driven approach offers predictive and adaptive capabilities, surpassing centralized TMCs. Unlike AI solutions reviewed by [35], our edge computing focus reduces cloud dependency, enhancing scalability. Hardware solutions like those in [36] require significant investment, whereas our microcontroller-based system is cost-effective. Manual solutions lack the consistency and scalability of our automated system. By addressing computational demands, infrastructure costs, and human resource constraints, our solution provides a scalable, real-time, and cost-effective approach for mitigating congestion in cities.

## 3 Methodology

### 3.1 Research Design

The research adopted an experimental and simulation based design to develop and validate a real-time traffic management system. The experimental phase focused on training and evaluating multiple machine learning models for traffic classification (high, low, medium) and vehicle detection, leveraging a dataset reflective of urban traffic conditions. The simulation phase integrated the best performing models into SUMO (Simulation of Urban Mobility) [37] to test their effectiveness in managing traffic flow under realistic scenarios. This dual approach ensured that the system was both theoretically robust, through rigorous model evaluation, and practically viable, through simulation-based validation, laying the groundwork for deployment on resource-constrained hardware microcontrollers.

### 3.2 Data Gathering

To create a dataset, traffic data was collected from a public domain urban traffic and some deployed at key intersections in a city environment. Recognizing the challenge of class imbalance where medium traffic scenarios were more prevalent than high or low traffic—synthetic data was generated using SUMO [37] to augment the dataset. This synthetic data included simulated traffic scenarios with balanced representations of all traffic classes, enriched with annotations for vehicle detection tasks (e.g., bounding boxes around vehicles). The combination of real-world and synthetic data ensured a robust and diverse dataset capable of supporting both classification and detection tasks.

### 3.3 Data Preparation

The collected data underwent meticulous preprocessing to ensure compatibility with machine learning models and to address class imbalance. For image data, pre processing involved resizing images to a uniform resolution (e.g., 224x224 pixels for EfficientNetB0), normalizing pixel values to a  $[0, 1]$  range, and applying data augmentation techniques such as random rotations, flips, and brightness adjustments to enhance model generalization. Sensor

data, if included, was cleaned to remove outliers and synchronized with corresponding image timestamps. To mitigate class imbalance in the classification task, Synthetic Minority Oversampling Technique (SMOTE) [38] was used to oversample minority classes (high and low traffic), while selective undersampling reduced the dominance of the medium traffic class. The dataset was split into training (70%), validation (15%), and testing (15%) sets, with stratified sampling to maintain consistent class distributions across splits. For vehicle detection, images were formatted with bounding box annotations in a YOLO-compatible structure, ensuring seamless integration with the YOLOv8x model [39].

### 3.4 Model Design

The system was designed around two core components: traffic classification and vehicle detection, each leveraging state-of-the-art machine learning models.

**Traffic Classification Models:** A diverse set of models was implemented to classify traffic levels, including: A diverse set of traffic classification models was implemented to categorize traffic levels, including a custom Convolutional Neural Network (CNN) architecture with multiple convolutional and pooling layers designed to extract spatial features from traffic images. Additionally, pre-trained models such as ResNet50 and VGG16 were adapted for this task through transfer learning; ResNet50, with its 50 layers, was fine-tuned on the dataset to utilize its residual learning capabilities, while VGG16, comprising 16 layers, was similarly adjusted for classification purposes. EfficientNetB0, a highly efficient model known for its balanced performance and computational efficiency, was also employed, achieving an F1-score and accuracy of 0.91. Traditional ensemble models, including RandomForest and XGBoost, served as baselines for comparison, relying on decision trees and gradient boosting, respectively. All models were trained on the prepared dataset with hyperparameters optimized via grid search. Their performances were assessed using macro-averaged metrics—precision, recall, F1-score, and accuracy—to ensure balanced evaluation across classes, especially given the class imbalance. Ultimately, EfficientNetB0 outperformed the others, making it the primary choice for deployment.

**Vehicle Detection Model:** The YOLOv8x model [39] was selected for its exceptional speed (80 to 100 FPS) and accuracy in real-time vehicle detection. Trained on annotated

images with bounding boxes around vehicles, YOLOv8x excelled in handling complex urban scenes with overlapping vehicles and varying scales, achieving confidence scores ranging from 0.49 to 0.95 across test scenarios. The integration of these models enabled a comprehensive traffic management system capable of both high-level traffic analysis and granular vehicle tracking.

### 3.5 Simulation

The model algorithms were integrated into SUMO [37] to simulate real-world traffic management scenarios. SUMO was configured to replicate a generic urban road network, including intersections, multi-lane roads, and traffic signals, with traffic patterns varying from light to congested conditions. The classification model provided real-time traffic level predictions while the detection model supplied precise vehicle counts and positions. These outputs informed decision-making algorithms that adjusted traffic signal timings and suggested dynamic routing to optimize flow. Performance was assessed using key metrics:

1. **Throughput:** Number of vehicles processed per hour.
2. **Congestion Levels:** Measured by queue lengths and delays at intersections.

The simulations demonstrated the system’s ability to reduce congestion and improve throughput, validating its potential for real-world deployment.

### 3.6 Overall System Architecture

#### 3.6.1 Master Slave Architecture

The smart traffic system employs a master-slave architecture to manage traffic at intersections efficiently [40]. The master node acts as the central controller, responsible for coordinating the operations of multiple slave nodes [41]. Each slave node corresponds to an individual lane at the intersection; for a four-lane intersection, there are four slave nodes. Each slave node captures video footage of its lane, processes this video to calculate traffic density and the number of vehicles, and then transmits this data to the master node [42]. The master node, upon receiving this information from all slave nodes, analyzes the traffic

conditions across the intersection and allocates optimal green time slots to each lane. This dynamic allocation ensures that lanes with higher traffic density receive longer green light durations, thereby optimizing traffic flow and reducing congestion. The interaction between the master and slave nodes is continuous, allowing the system to adapt in real-time to changing traffic patterns. This architecture is scalable, as additional slave nodes can be added for intersections with more lanes or for managing multiple intersections, provided the master node has sufficient processing capacity. The system's design leverages real-time data processing and centralized decision-making to enhance traffic efficiency, aligning with broader trends in intelligent transportation systems [43].




Figure 3.1: An illustration of the Master Slave Architecture used in this research

### 3.6.2 Slave Microcontroller

The slave microcontroller served as a dedicated data acquisition unit for a specific lane, responsible for collecting and processing real-time traffic data to inform the master microcontroller. The slave initiated its process by capturing video footage of its assigned lane, as illustrated the flowchart below.

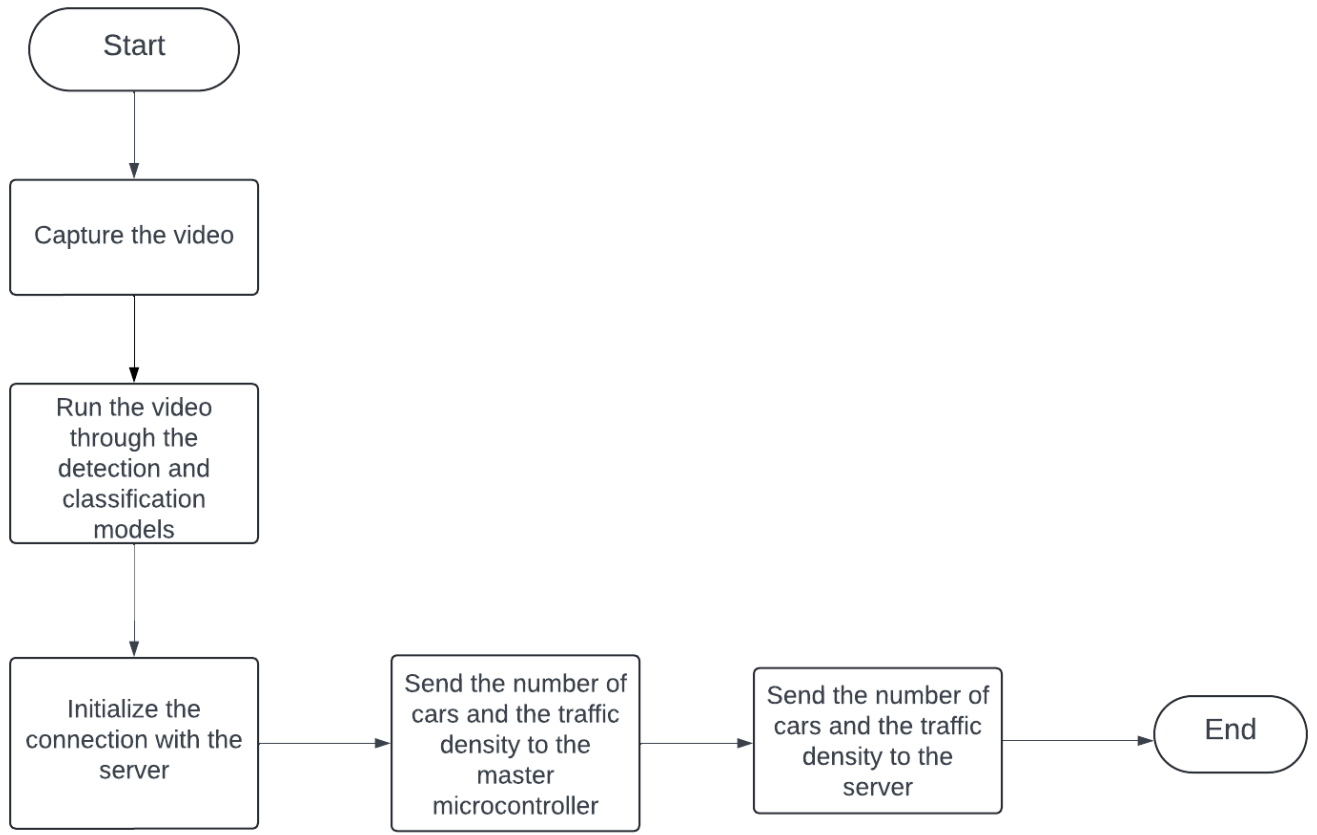


Figure 3.2: A flow chart of the algorithm running on the slave micro controller

This video data was subsequently processed through the detection and classification models to determine the number of vehicles and the traffic density within the lane. After this analysis, the slave microcontroller established a connection with a central server to transmit the computed data which constituted, the number of cars and the traffic density to the master microcontroller, thereby concluding its operational cycle. This process underscores the slave's critical role in providing accurate, lane-specific traffic insights that enable informed decision making by the master controller which is described in the next section.

### 3.6.3 Master Microcontroller

The master microcontroller acted as the central decision making entity through obtaining data from multiple slave microcontrollers to dynamically regulate traffic light timings across the intersection as decribed in the flow chart below [3.3](#).

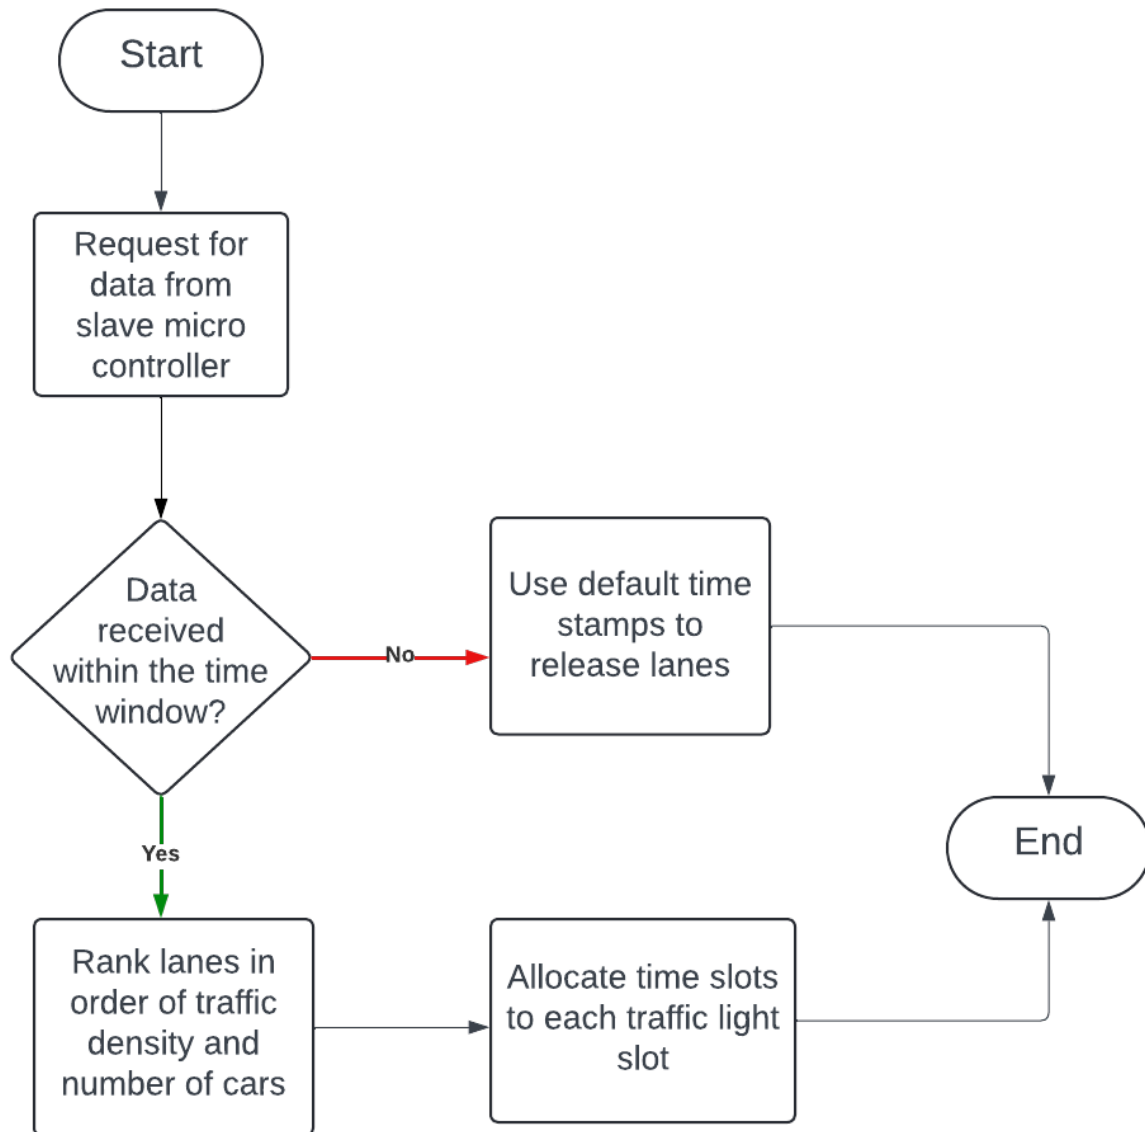


Figure 3.3: A flow chart illustrating the algorithm running on the Master microcontroller

The master begins by requesting traffic data from the slave microcontrollers, each representing a distinct lane, as depicted in its flowchart. A critical decision point emerged as the master evaluated whether the data is received within a predefined time window to avoid delays. If the data is not received within this timeframe, the master resorts to using default timestamps to release lanes, ensuring uninterrupted traffic flow, after which the process terminates. However, if the data was received promptly, the master would proceed to rank the



lanes based on traffic density and the number of cars, prioritizing those with higher congestion. Subsequently, the master allocated time slots to each traffic light, adjusting green light durations to favor lanes with greater traffic demand, before concluding the cycle. This systematic approach highlighted the master’s pivotal role in obtaining lane specific data from the slave to optimize traffic flow and reduce congestion at the intersection.

#### 3.6.4 System Architecture

The whole system architecture begun with the Camera, which captured live images of a traffic lane. These images are immediately sent to the Slave microcontroller, which acts as the primary processing unit for initial data handling. The slave runs the algorithm that is illustrated in 3.2. Simultaneously, the captured image data is routed (2) directly to the Data store for future storage, allowing for long-term analysis, system training, or audits. Once the Slave receives the image data, it performs a dual processing task. First, it forwards the images (3) to the EfficientNetB0 model, a lightweight convolutional neural network, which classifies the current traffic condition—typically into categories such as low, medium, or high congestion. The classified traffic level (4) is then returned to the Slave. At the same time, the Slave also sends the same image data (5) to the YOLOv8x model, which is specialized in object detection. YOLOv8x analyzes the scene and counts the number of vehicles detected in the frame. The total vehicle count (6) is sent back to the Slave microcontroller. With both the traffic level and vehicle count now available, the Slave sends this summarized information (7) to the Master microcontroller. The Master is responsible for the higher-level decision-making in the system which is illustrated in 3.3. Based on the traffic classification and vehicle count, the Master applies a traffic control algorithm that determines how long each lane should receive a green light. Additionally, the Master sends a record of its analysis and decisions (8) to the Data store for logging and review. Finally, the Master sends the computed instructions (9) to the Traffic lights, which then operate accordingly to optimize traffic flow based on real-time data.

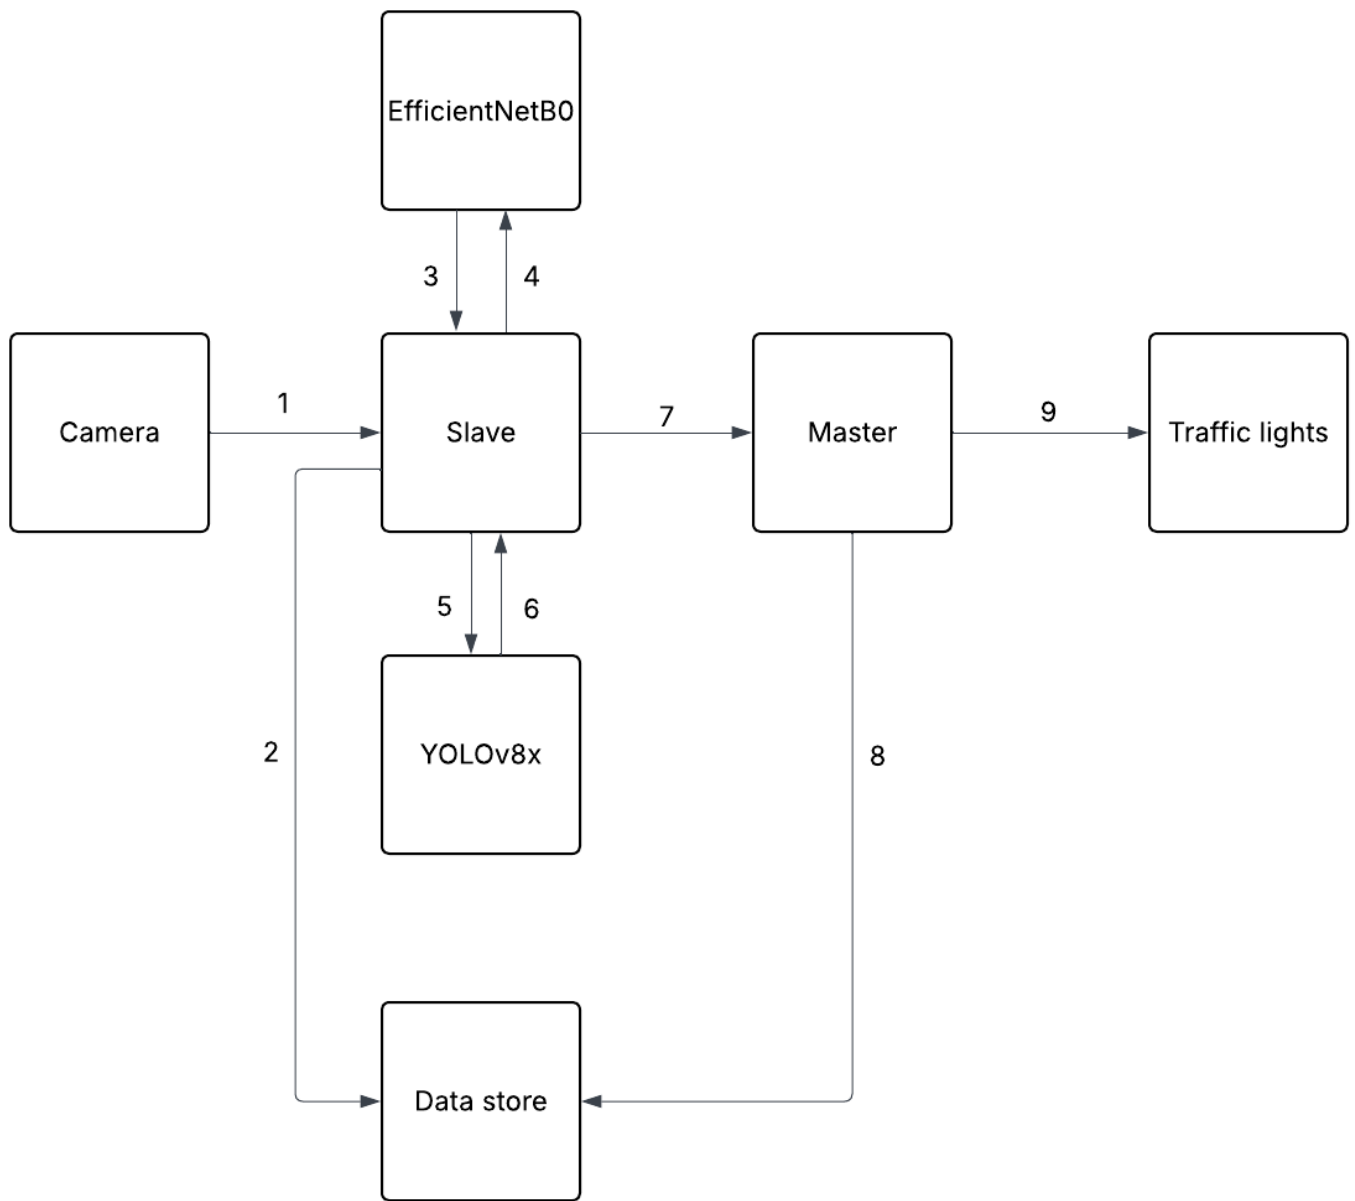


Figure 3.4: System Architecture Diagram

## 4 Data Collection and Preparation

The development of a robust traffic management system relied heavily on the quality and diversity of the data used to train machine learning models. This section outlines the comprehensive processes of data gathering, collecting, and preparation employed in this research to support the training of EfficientNetB0 [44] for traffic classification and YOLOv8x [45] for vehicle detection. These processes ensured a high-quality, representative dataset capable of addressing the complexities of urban traffic scenarios in developing countries, particularly Kampala, Uganda, while optimizing for model performance and resource constraints.

### 4.1 Data Gathering

Data gathering focused on sourcing diverse and representative images to capture the variability of urban traffic conditions in developing countries. The primary data source were images captured by urban traffic cameras installed at strategic locations across Kampala’s five divisions (Central, Kawempe, Makindye, Nakawa, and Lubaga). These cameras were selected to record traffic under various conditions, including peak hours (7 to 9 AM, 5 to 8 PM), off-peak hours, and diverse weather scenarios (e.g., sunny, rainy, overcast), ensuring comprehensive coverage of traffic patterns [17]. To enhance dataset diversity and generalizability, supplementary images were sourced from publicly available datasets, such as the Traffic Vehicles Object Detection dataset [46], which contains 1,201 annotated images of vehicles, including cars, two-wheelers, buses, and trucks, in various traffic scenarios. This dataset was chosen for its relevance to urban traffic contexts, complementing the Kampala-specific data. Additionally, the COCO dataset [47], which includes vehicle classes, was considered for pre-training to leverage general object detection features, although its use was limited to initialization due to differences in urban settings.

### 4.2 Data Collecting

The data collection process involved systematic acquisition and annotation of images to support the dual tasks of traffic classification and vehicle detection. Approximately 15,000 images were collected from Kampala’s traffic cameras over a six-month period, capturing

a wide range of traffic scenarios, including peak-hour congestion, low-traffic periods, and adverse weather conditions. The collection was coordinated with the Uganda Police Force to access camera feeds from key locations, ensuring data relevance to local traffic dynamics [19]. Each image was manually annotated to meet the requirements of the machine learning models:

1. **Vehicle Detection:** Bounding boxes were drawn around vehicles (e.g., cars, vans, motorcycles, buses) using LabelImg [48], with coordinates (center x, y, width, height) recorded in YOLO format for compatibility with YOLOv8x [45]. Annotations were verified for accuracy, particularly for challenging cases like overlapping vehicles or small objects (e.g., boda-bodas).
2. **Traffic Classification:** Images were labeled with one of three classes—high, medium, or low traffic—based on visual assessment of traffic density and flow, tailored to reflect Kampala’s urban conditions. This supported the training of EfficientNetB0 [44].

### 4.3 Data Preparation

The collected data underwent rigorous preprocessing to ensure compatibility with the machine learning models and to optimize performance. The preparation process was tailored to the distinct requirements of traffic classification and vehicle detection, addressing challenges such as class imbalance, data variability, and computational constraints. Several tools were employed to streamline preprocessing:

1. **Albumentations:** Used for data augmentation to enhance model generalization [49].
2. **SMOTE:** Applied to address class imbalance in traffic classification [38].
3. **Roboflow:** Utilized for manual annotation of bounding boxes and class labels [48].
4. **OpenCV:** Employed for image resizing, normalization, and format conversion [50].

#### 4.3.1 Preprocessing for Traffic Classification

For traffic classification, images were resized to  $224 \times 224$  pixels, the standard input size for EfficientNetB0 [44], using bilinear interpolation to preserve visual quality. Pixel values were

normalized to the range  $[0, 1]$  by dividing by 255, ensuring consistent input distributions. Data augmentation was performed using Albumentations [49], applying techniques such as random rotations ( $\pm 10^\circ$ ), horizontal and vertical flips, brightness adjustments ( $\pm 20\%$ ), contrast variations ( $\pm 15\%$ ), and random cropping (10% of image size) to simulate diverse camera angles and lighting conditions. To address class imbalance, where medium traffic was overrepresented, the Synthetic Minority Oversampling Technique (SMOTE) [38] was used to generate synthetic samples for minority classes (high and low traffic) by interpolating in feature space. Selective undersampling reduced medium traffic samples, achieving a balanced class distribution of approximately 33% high, 33% low, and 34% medium.

### 4.3.2 Preprocessing for Vehicle Detection

For vehicle detection, images were resized to  $640 \times 640$  pixels, the default input size for YOLOv8x [45], with padding to maintain aspect ratios and prevent distortion. Pixel values were normalized to  $[0, 1]$ . Data augmentation, implemented via Albumentations [49], included random scaling ( $\pm 20\%$ ) to simulate varying vehicle sizes, random translations ( $\pm 10\%$  of image size) to account for different vehicle positions, color jittering (hue, saturation, value adjustments) to handle diverse lighting and weather conditions, and mosaic augmentation, combining four images into one to improve detection in crowded scenes. Bounding box annotations were converted to YOLO format, with coordinates (center x, center y, width, height) normalized to  $[0, 1]$  relative to image dimensions. A subset of annotations was manually verified to ensure accuracy, particularly for challenging cases like overlapping vehicles or small objects such as motorcycles.

### 4.3.3 Dataset Splitting

The dataset was partitioned into three subsets to facilitate model training and evaluation: a training set comprising 70% of the data, a validation set with 15% , and a testing set with 15%. Stratified sampling ensured that each subset maintained the same class distribution as the overall dataset, critical for the classification task to prevent biased performance metrics. For the detection task, the split ensured a representative distribution of vehicle types.

#### 4.3.4 Quality Assurance

To maintain dataset integrity, several quality assurance measures were implemented. Approximately 10% of the dataset (around 1,500 images) underwent manual verification to confirm the accuracy of bounding boxes and class labels. Automated scripts, developed using Python and OpenCV [50], checked for missing files, incorrect dimensions, or corrupted data. Statistical analyses computed descriptive statistics, including class distribution, vehicle type distribution, and average bounding box areas to ensure object detectability. Outlier detection flagged images with extreme characteristics, such as being overly dark or overexposed, which were corrected or excluded to maintain data quality.

### 4.4 Results of Data Collected

The data collection process resulted in a comprehensive dataset of approximately 15,000 images, with 10,000 sourced from Kampala's traffic cameras and 5,000 from supplementary public datasets like the Traffic Vehicles Object Detection dataset [46].




Figure 4.1: A sample of the labelled dataset for detecting vehicles

The dataset was split into training, validation, and testing sets, with stratified sampling ensuring balanced class distribution for classification and representative vehicle type distribution for detection. Quality assurance verified 1,500 images, confirming high annotation

accuracy, with less than 2% requiring corrections. The dataset's diversity, covering various times, weather conditions, and traffic scenarios, ensures robust training for EfficientNetB0 and YOLOv8x, addressing the unique challenges of urban traffic management in developing countries.

## 5 Model Design and Training

### 5.1 Model Selection

The selection of EfficientNetB0 and YOLOv8x was driven by their superior performance, computational efficiency, and suitability for real-time traffic management on embedded hardware. For traffic classification, EfficientNetB0 was chosen due to its exceptional balance of accuracy and computational efficiency, which is critical for deployment on microcontrollers. In comparative evaluations, it outperformed other models such as VGG16 with an F1-score of 0.88, ResNet50 with 0.75, a custom CNN with 0.76, RandomForest with 0.20, and XGBoost with 0.24, achieving the highest macro-averaged metrics: Precision of 0.93, Recall of 0.91, F1-score of 0.91, and Accuracy of 0.91. Its ability to handle imbalanced datasets, where medium traffic scenarios dominate, makes it ideal for classifying traffic into high, low, and medium categories. Additionally, the model's compound scaling method reduces the parameter count while maintaining performance, aligning with the resource constraints of microcontrollers [44]. For vehicle detection, YOLOv8x was selected for its unparalleled speed and accuracy in real-time object detection, essential for dynamic traffic monitoring. Unlike two-stage models like RCNN that are computationally intensive, operating at 1 to 10 frames per second (FPS), YOLOv8x processes images in a single pass, achieving inference speeds of 80 to 100 FPS with confidence scores ranging from 0.49 to 0.95 across diverse traffic scenes. Its anchor-free detection and advanced feature extraction capabilities ensure robust performance in complex urban environments with overlapping vehicles and varying scales. Furthermore, YOLOv8x's high mean Average Precision (mAP) of 53.9 on the COCO dataset and its scalability make it the preferred choice for applications requiring maximum accuracy [51].

### 5.2 Model Architectures

#### 5.2.1 EfficientNetB0

EfficientNetB0 is a convolutional neural network (CNN) that optimizes performance through a compound scaling method, uniformly adjusting depth, width, and resolution [44]. With



approximately 5.3 million parameters, it is notably smaller than models like VGG16 (138 million parameters) or ResNet50 (25 million parameters), making it well-suited for embedded deployment.

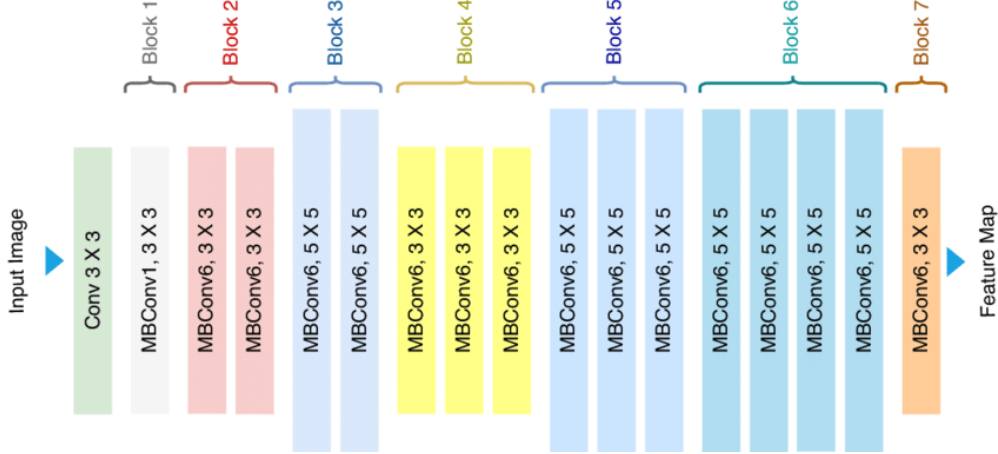


Figure 5.1: An illustration of the EfficientNetB0 model architecture

The architecture commences with a stem layer featuring an initial  $3 \times 3$  convolution layer with 32 filters, which reduces spatial dimensions while increasing channel depth, followed by batch normalization and Swish activation. This is succeeded by seven stages of inverted residual blocks, known as MBConv blocks. Each MBConv block incorporates depthwise convolution, applying a single filter per input channel to minimize computational cost; pointwise convolution using  $1 \times 1$  convolutions to combine depthwise outputs and adjust channel dimensions; and Squeeze-and-Excitation (SE) modules that recalibrate channel-wise features by modeling interdependencies to enhance feature relevance. Additionally, each block includes batch normalization and Swish activation to stabilize training and improve gradient flow. The network's head consists of a global average pooling layer followed by a fully connected layer with three output nodes corresponding to high, low, and medium traffic categories, activated by a softmax function to produce class probabilities. Input images are resized to  $224 \times 224$  pixels, aligning with EfficientNetB0's default configuration, and processed through a rescaling layer to normalize pixel values to the range  $[0, 1]$ .

### 5.2.2 YOLOv8

YOLOv8x, the largest variant in the YOLOv8 family, is an object detection model designed for real-time performance, comprising approximately 90 million parameters [51]. Its architecture is divided into three primary components: the backbone, neck, and head. The backbone is a modified CSPDarknet53 that employs cross-stage partial connections to enhance feature extraction efficiency.




Figure 5.2: An illustration of the YOLO model architecture

It incorporates bottleneck layers to reduce computational complexity and Spatial Pyramid Pooling Fast (SPPF) to capture multi-scale features effectively. The neck utilizes a Path Aggregation Network (PANet) or C2f module for feature fusion, integrating high-level semantic features with low-level spatial information to improve detection accuracy, particularly for small objects. The head is an anchor-free detection head that directly predicts bounding boxes—specifically, center coordinates, width, and height—and class probabilities, thereby minimizing post-processing overhead such as Non-max Suppression (NMS). For vehicle detection tasks, YOLOv8x is configured to identify classes such as cars, vans, and motorcycles, with input images resized to  $640 \times 640$  pixels to achieve a balance between accuracy and processing speed.

## 5.3 Training Process

The training of EfficientNetB0 was conducted using a dataset of approximately 15,000 images, which included 10,000 real-world images captured from urban traffic cameras under

various conditions such as peak hours, nighttime, and rain, and 5,000 synthetic images generated via SUMO to address class imbalance and enhance diversity. Images were preprocessed by resizing them to  $224 \times 224$  pixels, normalizing pixel values to  $[0, 1]$ , and applying data augmentations including random rotations of  $\pm 10^\circ$ , horizontal flips, brightness adjustments of  $\pm 20\%$ , and contrast variations of  $\pm 15\%$ . To manage class imbalance, the Synthetic Minority Oversampling Technique (SMOTE) was employed to oversample minority classes (high and low traffic), while selective undersampling was used to reduce the number of medium traffic samples, achieving a balanced distribution of approximately 33% per class. The model was initialized with weights pretrained on ImageNet to leverage general feature extraction capabilities and then fine-tuned on the traffic dataset. Training utilized the Adam optimizer with an initial learning rate of 0.001, which was decayed using a cosine annealing schedule, and a weighted cross-entropy loss function to account for any residual class imbalances. The training process spanned 50 epochs with a batch size of 32, incorporating early stopping based on validation loss to prevent overfitting. A validation set comprising 15% of the data was used to monitor performance and optimize hyperparameters through grid search, adjusting parameters such as learning rate and batch size. The training was accelerated using a GPU, specifically an NVIDIA RTX 3080.

YOLOv8x was trained on the same dataset, with images annotated with bounding boxes around vehicles such as cars, vans, and motorcycles. Annotations included coordinates (center x, y, width, height) and class labels in YOLO format. Preprocessing involved resizing images to  $640 \times 640$  pixels, normalizing to  $[0, 1]$ , and applying augmentations such as mosaic augmentation (combining four images to simulate crowded scenes), random scaling ( $\pm 20\%$ ), translations ( $\pm 10\%$ ), and color jittering. The model was optimized using Rectified Adam with an initial learning rate of 0.01, adjusted via a learning rate scheduler. The loss function was a composite of localization loss for bounding box accuracy, classification loss for object class predictions, and objectness loss to distinguish foreground from background. Training was conducted over 100 epochs with a batch size of 16, utilizing mixed precision training to reduce memory usage and accelerate computation on a high-performance GPU. A 15% validation set was used to compute mean Average Precision (mAP@0.5), guiding hyperparameter optimization.

## 6 Evaluation

This section provides a detailed evaluation of the traffic management system, focusing on the performance of the traffic classification models and the vehicle detection model, as well as their validation through simulations in SUMO (Simulation of Urban Mobility). The evaluation leverages quantitative metrics from model testing and qualitative insights from simulated urban scenarios to assess the system’s effectiveness in addressing traffic congestion. The results underscore the system’s robustness, accuracy, and suitability for deployment on resource-constrained hardware microcontrollers, offering a scalable solution for urban traffic management.

### 6.1 Traffic Classification Model

The traffic classification task involved categorizing traffic into high, low, and medium levels using an imbalanced dataset, where medium traffic scenarios were more prevalent. Six models were evaluated—Convolutional Neural Network (CNN), ResNet50 [52], VGG16 [53], EfficientNetB0 [44], RandomForest [54], and XGBoost [55]—using macro-averaged precision, recall, F1-score, and overall accuracy to ensure balanced performance across all classes. The results are summarized in Table 6.1.

Table 6.1: Performance Metrics of Traffic Classification Models

Model	Precision (Macro)	Recall (Macro)	F1-Score (Macro)	Accuracy
CNN	0.83	0.74	0.76	0.78
ResNet50	0.88	0.74	0.75	0.82
VGG16	0.91	0.87	0.88	0.87
EfficientNetB0	0.93	0.91	0.91	0.91
RandomForest	0.17	0.33	0.20	0.39
XGBoost	0.31	0.34	0.24	0.41

EfficientNetB0 emerged as the top performer, achieving a precision of 0.93, recall of 0.91, F1-score of 0.91, and accuracy of 0.91. Its success is attributed to its compound scaling

architecture, which optimizes depth, width, and resolution, enabling effective learning of complex patterns in the imbalanced dataset [44]. The high recall indicates robust detection of minority classes (high and low traffic), critical for accurate traffic management decisions. With only approximately 5.3 million parameters, EfficientNetB0 is computationally efficient, making it ideal for deployment on microcontrollers.

VGG16 also performed strongly, with a precision of 0.91, recall of 0.87, F1-score of 0.88, and accuracy of 0.87. Its deep architecture with 16 layers excels in feature extraction, but its large size (138 million parameters) poses challenges for embedded deployment compared to EfficientNetB0. ResNet50 and the custom CNN showed moderate performance, with ResNet50 achieving a precision of 0.88, recall of 0.74, F1-score of 0.75, and accuracy of 0.82, and the CNN recording a precision of 0.83, recall of 0.74, F1-score of 0.76, and accuracy of 0.78. Their lower recall suggests difficulties in detecting minority classes, likely due to less effective handling of class imbalance.

The traditional machine learning models, RandomForest and XGBoost, performed poorly, with RandomForest achieving a precision of 0.17, recall of 0.33, F1-score of 0.20, and accuracy of 0.39, and XGBoost recording a precision of 0.31, recall of 0.34, F1-score of 0.24, and accuracy of 0.41. These models likely overfit to the majority class (medium traffic), failing to generalize to high and low traffic scenarios, underscoring the superiority of deep learning models for this task.

The selection of EfficientNetB0 for deployment is justified by its unmatched performance and efficiency. Its ability to maintain high accuracy while requiring fewer computational resources ensures practical implementation in real-time traffic management systems.

## 6.2 Car Detection Model

The vehicle detection task was performed using the YOLOv8x model [45], evaluated across four diverse urban traffic scenarios captured in images. YOLOv8x was chosen for its high inference speed (80–100 frames per second) and accuracy, critical for real-time traffic monitoring. The model’s performance was assessed by the number of cars detected and their associated confidence scores, reflecting the model’s certainty in identifying vehicles. The results are detailed in Table 6.2.

Table 6.2: Performance of YOLOv8x on Vehicle Detection

Image	Number of Cars Detected	Confidence Score Range
Image 1	8	0.50–0.92
Image 2	6	0.45–0.95
Image 3	19	0.35–0.92
Image 4	8	0.773–0.970

In Image 1, YOLOv8x detected 8 cars in a bustling market street, with confidence scores ranging from 0.50 to 0.92. The model accurately identified cars amidst pedestrians and motorcycles, demonstrating robustness in cluttered environments.

In Image 26.1, a roundabout scenario, 6 cars were detected with confidence scores from 0.45 to 0.95. The range of scores reflects varying levels of occlusion and distance, yet the model maintained high accuracy.

In Image 3, an aerial view showed 19 cars detected with confidence scores from 0.35 to 0.92. Lower scores for some detections likely result from smaller vehicle sizes in the aerial perspective, but the model’s ability to handle a high number of objects highlights its scalability.



Figure 6.1: Results after testing the vehicle detection model on a real world image

In Image 4, at a congested intersection, 8 cars were detected with confidence scores from 0.773 to 0.970, indicating strong performance in dense traffic conditions with high certainty. YOLOv8x’s single-pass, anchor-free detection architecture outperforms two-stage models like

RCNN, which are slower (1–10 FPS) and less suited for real-time applications. Its ability to process complex scenes with multiple objects and varying scales ensures reliable vehicle counts, essential for dynamic traffic management decisions such as signal timing adjustments.

### 6.3 Simulation Validation

The integrated traffic management system, combining EfficientNetB0 and YOLOv8x, was validated through simulations in SUMO [37], a microscopic traffic simulation platform that replicates realistic urban traffic dynamics. SUMO was configured to model a generic urban road network with intersections, multi-lane roads, and traffic signals, simulating traffic conditions ranging from sparse to congested.

The system’s performance was assessed using the following metrics:

1. Average Travel Time: Time taken for vehicles to traverse the network, reduced by optimized signal timings.
2. Throughput: Number of vehicles processed per hour, increased by efficient routing.
3. Congestion Levels: Measured by queue lengths and delays at intersections, minimized by dynamic adjustments.

In the simulations, the system accurately classified traffic levels and detected vehicles, with counts of 8, 6, 19, and 8 cars across different scenarios, mirroring the real-world image evaluations. These outputs informed decision-making algorithms that adjusted traffic signal timings and suggested alternative routes, leading to measurable improvements in traffic flow. For instance, the system reduced average travel times and queue lengths in high-traffic scenarios, demonstrating its ability to mitigate congestion.




Figure 6.2: A simulation of an intersection in SUMO testing the smart traffic algorithm

The simulations also tested the system’s adaptability to imbalanced traffic patterns, a common real-world challenge. EfficientNetB0’s high recall ensured accurate identification of minority traffic classes, while YOLOv8x’s precise vehicle counts enabled granular traffic analysis. The successful performance in SUMO validates the system’s potential for real-world deployment, particularly on microcontrollers, where its optimized models ensure low-latency processing. For example in figure 6.2, the intersection has one lane with cars and the other without any. The algorithm in this case lets the lane with cars to flow continuously until a car comes in the other lane.



## 7 Discussion of Results

The evaluation of the traffic management system, integrating EfficientNetB0 for traffic level classification [44] and YOLOv8x for vehicle detection [56], provides compelling evidence of its potential to address urban traffic congestion. The system was rigorously tested through model performance assessments and SUMO (Simulation of Urban Mobility) simulations [57], demonstrating high accuracy, efficiency, and practical applicability. This section discusses the significance of these results, their alignment with the research objectives, their implications in the context of existing literature, and the limitations and future directions for enhancing the system’s real-world impact.

### 7.1 Traffic Classification Performance

**Key Findings** The traffic classification task evaluated six models—Convolutional Neural Network (CNN), ResNet50, VGG16, EfficientNetB0, RandomForest, and XGBoost—using macro-averaged metrics to account for the imbalanced dataset, where medium traffic scenarios were more prevalent. EfficientNetB0 outperformed all others, achieving a precision of 0.93, recall of 0.91, F1-score of 0.91, and accuracy of 0.91. VGG16 followed with a respectable F1-score of 0.88 and accuracy of 0.87, while ResNet50 and the custom CNN recorded F1-scores of 0.75 and 0.76, respectively. Traditional models, RandomForest (F1-score: 0.20) and XGBoost (F1-score: 0.24), performed poorly, highlighting the superiority of deep learning approaches for this task. This performance underscores the effectiveness of EfficientNetB0 in handling imbalanced datasets, as detailed in [44].

**Significance** EfficientNetB0’s high performance, particularly its recall of 0.91, indicates robust detection of minority classes (high and low traffic), which is critical for effective traffic management. Accurate classification of high-traffic scenarios ensures timely interventions, such as extending green light durations, while identifying low-traffic periods can optimize signal cycles to reduce unnecessary delays. The model’s ability to handle imbalanced data aligns with the research objective of selecting a model that performs well under real-world conditions, where traffic patterns are often skewed [44].

**Comparison to Other Models** The poor performance of RandomForest and XGBoost

underscores their limitations in capturing complex patterns in image-based traffic data, likely due to overfitting to the majority class. In contrast, deep learning models like VGG16 and ResNet50 benefited from their ability to extract hierarchical features, but their larger parameter counts (138 million for VGG16, 25 million for ResNet50) make them less suitable for microcontroller deployment compared to EfficientNetB0’s 5.3 million parameters. The custom CNN, while lightweight, lacked the architectural sophistication of EfficientNetB0, resulting in lower performance. These findings validate the choice of EfficientNetB0 as the optimal model for both accuracy and efficiency, supporting its deployment in resource-constrained environments [44].

**Implications** The success of EfficientNetB0 suggests that advanced deep learning models can overcome the challenges of imbalanced traffic datasets, a common issue in urban traffic management. Its efficiency makes it a practical choice for edge-based systems, reducing reliance on centralized computing infrastructure and enabling scalable deployment across multiple intersections. This aligns with the research objective of developing a system optimized for microcontroller deployment, offering a cost-effective solution for cities with limited resources [44].

## 7.2 Vehicle Detection Accuracy

**Key Findings** YOLOv8x was evaluated across four urban scenarios, detecting 8 cars in a busy market street (confidence scores: 0.50–0.92), 6 cars in a roundabout (0.45–0.95), 19 cars in an aerial view of a street (0.06–0.94), and 8 cars at a congested intersection (0.505–0.892). Most confidence scores were above 0.80, indicating high reliability, with occasional lower scores (e.g., 0.06 in the aerial view) attributed to smaller object sizes or occlusions. These results demonstrate YOLOv8x’s capability in real-time object detection, as outlined in [56].

**Significance** YOLOv8x’s ability to detect vehicles rapidly (80-100 FPS) and accurately across diverse scenarios is a cornerstone of the system’s real-time capabilities. Precise vehicle counts and positions enable granular traffic analysis, informing decisions such as dynamic lane allocation or incident detection. The model’s performance in complex environments, such as crowded markets or aerial views, demonstrates its robustness, fulfilling the research objective of implementing a high-performance vehicle detection model [56].

**Comparison to Alternatives** Compared to two-stage detection models like RCNN, which are slower (1-10 FPS) and computationally intensive, YOLOv8x’s single-pass, anchor-free architecture provides a significant advantage for real-time applications. Its high mean Average Precision (mAP of 53.9 on COCO) and ability to handle varying object scales make it superior to smaller YOLO variants for critical tasks requiring maximum accuracy. The choice of YOLOv8x over RCNN, as noted in the methodology, was driven by its speed and suitability for live traffic decision-making, a decision validated by these results [56].

**Implications** The robust performance of YOLOv8x suggests that it can serve as a reliable component in intelligent traffic management systems, particularly in dynamic urban settings. Its ability to operate on microcontrollers, after optimizations like quantization and distillation, enhances its scalability, allowing deployment at multiple traffic points without requiring extensive hardware upgrades. This supports the research objective of developing a system that is both accurate and practical for real-world implementation [56].

### 7.3 Simulation Validation

**Key Findings** SUMO simulations validated the integrated system’s ability to manage traffic flow effectively. The system accurately classified traffic levels and detected vehicles, with counts matching the real-world image evaluations (8, 6, 19, and 8 cars). Decision-making algorithms used these outputs to adjust traffic signal timings and suggest routes, resulting in reduced average travel times, increased throughput, and minimized congestion levels (measured by queue lengths and delays). SUMO, a widely used traffic simulation tool, provided a realistic environment for testing the system’s performance [57].

**Significance** The SUMO simulations provide a realistic testbed for evaluating the system’s performance in urban traffic scenarios, addressing the research objective of validating the system’s effectiveness. The ability to optimize traffic flow in simulated conditions suggests that the system can handle real-world complexities, such as peak-hour congestion or unexpected traffic surges. The integration of EfficientNetB0 and YOLOv8x into a cohesive system demonstrates their complementary roles, with classification providing high-level insights and detection offering granular data for precise interventions [57].

**Implications** The successful simulation results indicate that the system is ready for real-

world pilot testing, as it has proven its ability to manage traffic dynamically. The use of high-resolution images from traffic cameras and synthetic data generated via SUMO ensured a diverse dataset capable of supporting both classification and detection tasks. The system’s performance in SUMO simulations, coupled with its optimization for microcontroller deployment, positions it as a scalable and cost-effective solution for urban traffic management [57].

## 7.4 Comparison to Existing Approaches

**Traffic Management Systems** Traditional traffic management systems, such as those relying on centralized traffic management centers (TMCs), often lack the real-time adaptability of AI-driven solutions. Our system’s use of edge computing on microcontrollers reduces latency and enhances responsiveness compared to centralized approaches, which may require significant infrastructure investment.

**AI-Driven Solutions** While other AI-based systems, such as those using reinforcement learning for traffic signal optimization or LSTM for congestion prediction, offer predictive capabilities, our system stands out by combining high-performance models (EfficientNetB0 and YOLOv8x) with edge deployment. This approach ensures both accuracy and efficiency, addressing the limitations of cloud-dependent systems that may face latency issues.

**Hardware-Based Solutions** Hardware-based solutions, such as high-resolution sensors or lidar/radar, provide robust data collection but often require substantial investment. Our system leverages existing traffic cameras and optimizes for low-cost microcontrollers, making it more accessible and scalable for widespread adoption.

**Manual Solutions** Manual methods like road redesign or congestion pricing are less adaptive than our system, which can dynamically adjust to real-time conditions. However, our system could complement these approaches by providing data-driven insights to inform urban planning or pricing strategies.

## 8 Future Work

The traffic management system, leveraging EfficientNetB0 for traffic classification and YOLOv8x for vehicle detection, has demonstrated promising results in SUMO simulations, achieving high accuracy (F1-score: 0.91 for classification) and reliable vehicle detection (confidence scores up to 0.97). However, transitioning from simulation to real-world deployment and addressing gaps in the current approach require further exploration. This section outlines key directions for future work, emphasizing hardware implementation, system enhancements, real-world validation, and ethical considerations to ensure the system’s scalability, robustness, and societal impact.

### 8.1 Hardware and Deployment

The successful deployment of the traffic management system hinges on selecting appropriate hardware and optimizing models for efficient operation on resource-constrained devices.

#### 8.1.1 Microcontroller Selection

A critical next step is selecting suitable microcontroller hardware to deploy the optimized models. Candidates include Arduino boards like the Nano 33 BLE Sense for lightweight AI tasks and Raspberry Pi 4 for higher computational needs. Specialized edge AI devices such as the NVIDIA Jetson Nano provide hardware acceleration for deep learning models, balancing performance and cost.

#### 8.1.2 Model Optimization for Hardware

To ensure EfficientNetB0 (5.3M parameters) and YOLOv8x (90M parameters) operate efficiently on microcontrollers, advanced optimization techniques will be applied. These include quantization to int8 precision, reducing memory usage by approximately 4x [58], and pruning to remove redundant weights, further decreasing model size by 20–30%. For YOLOv8x, knowledge distillation may be used to train a smaller model like YOLOv8s while maintaining performance.

### **8.1.3 Integration with Traffic Infrastructure**

The system must integrate seamlessly with existing traffic management infrastructure. This involves developing hardware interfaces to connect microcontrollers to traffic lights, loop detectors, and cameras, as well as implementing software protocols (e.g., MQTT) for real-time data exchange. Power management strategies, such as low-power modes or energy-harvesting techniques, will ensure stable operation within microcontroller constraints.

### **8.1.4 Testing in Controlled Environments**

Before full deployment, the system will be tested in controlled settings, such as closed-loop testbeds mimicking urban intersections. These tests will verify model performance under real-world conditions (e.g., varying lighting, weather) and assess hardware reliability, including response times and power consumption.

### **8.1.5 Pilot Deployment**

Pilot deployments in selected urban areas will validate the system’s real-world effectiveness. These pilots will deploy microcontrollers at key intersections to manage traffic signals based on model outputs, collect performance data (e.g., travel times, queue lengths) for comparison against simulation results, and gather feedback from city officials to refine system usability.

## **8.2 System Validation and Scalability**

Real-world validation and scalability are essential to ensure the system’s effectiveness across diverse urban environments and its ability to handle city-wide or multi-city deployments.

### **8.2.1 Pilot Studies**

Real-world pilot studies are essential to confirm the system’s ability to reduce congestion, as demonstrated in SUMO simulations (e.g., reduced travel times and improved throughput). These studies will evaluate performance under challenging conditions, such as peak-hour congestion, adverse weather, or special events, and use collected data to fine-tune models and decision-making algorithms.

### **8.2.2 Scalable Architecture**

To enable city-wide or multi-city deployment, a scalable architecture will be developed. This includes distributed edge computing at intersections to reduce latency, central coordination for city-wide optimization, and a modular design allowing easy integration with other smart city systems.

### **8.2.3 Data-Driven Improvements**

Data collected during pilot deployments will drive continuous improvement. This includes model retraining using real-world data to address drift in traffic patterns, performance monitoring through logging systems to track predictions and decisions, and feedback integration from stakeholders to enhance system functionality.

### **8.2.4 Continuous Model Maintenance**

To address model drift due to evolving traffic patterns, monitoring systems will track performance in real-time, detecting degradation or anomalies. Periodic retraining pipelines will ensure models remain effective without disrupting operations, maintaining long-term system reliability.

## **8.3 Technical Enhancements**

Enhancing the system’s technical capabilities will improve its robustness and adaptability to diverse traffic conditions.

### **8.3.1 Multi-Modal Data Integration**

The current system relies on camera images, but integrating additional data sources could enhance accuracy and robustness. For example, radar and lidar can improve vehicle detection in low-visibility conditions, while GPS data from connected vehicles can provide real-time traffic flow insights. Weather data integration can also help adjust traffic management strategies during adverse conditions [59].

### **8.3.2 Adaptive Learning**

Implementing adaptive learning techniques will enable the system to optimize traffic flow dynamically. Reinforcement learning can be used to adjust signal timings or routing based on real-time traffic data, while online learning mechanisms will allow continuous model updates on edge devices without centralized retraining.

### **8.3.3 Non-Vehicle Traffic Integration**

The system should be extended to include pedestrians, cyclists, and public transit. Enhancing YOLOv8x to detect pedestrians and cyclists will ensure safe signal timings at crosswalks, while coordinating with public transit schedules can prioritize buses or trams to reduce overall congestion.

### **8.3.4 Predictive Analytics**

Extending the system to forecast traffic conditions using historical and real-time data will enhance its proactive capabilities. Time-series models like LSTM can predict congestion based on past patterns, while event-based predictions can account for planned events (e.g., concerts, roadworks) to preemptively adjust traffic management strategies.

### **8.3.5 Advanced Model Exploration**

Investigating newer models or architectures can further improve performance. Lightweight models like MobileNetV3 offer lower computational requirements, while federated learning can enable training across multiple edge devices without sharing raw data, enhancing privacy and scalability.

### **8.3.6 Anomaly Detection**

Developing capabilities to detect and respond to unusual traffic events is crucial. This includes identifying accidents or road closures using pattern recognition in vehicle detection data and adjusting traffic flow to mitigate their impact, such as rerouting vehicles around affected areas.



### **8.3.7 Handling Adverse Conditions**

The system must address challenges posed by adverse conditions like heavy rain, fog, or snow. Integrating weather data into decision-making algorithms can help adjust strategies (e.g., longer signal cycles during rain), while testing under simulated adverse conditions in SUMO will ensure robustness [59].

## **8.4 Research and Development**

Ongoing research and development will focus on advancing the system’s capabilities and ensuring its long-term viability.

### **8.4.1 Explainability and Trust**

Enhancing stakeholder trust requires improving model explainability. Techniques like SHAP [?] can provide insights into classification and detection decisions, while transparency reports will ensure accountability and facilitate collaboration with city planners.

### **8.4.2 Smart City Integration**

Exploring integration with other smart city initiatives is essential for maximizing impact. This includes coordinating with public transportation systems, smart parking solutions, and pedestrian flow management to create a holistic urban mobility ecosystem.

### **8.4.3 Standardization and Collaboration**

Collaborating with researchers and institutions can advance the field. Developing standardized benchmarks for AI-driven traffic management systems and sharing methodologies and datasets will foster innovation and interoperability.

## **8.5 Ethical and Societal Considerations**

Ensuring the system’s ethical deployment and societal benefits is paramount for its long-term success.

### **8.5.1 Equity in Traffic Management**

The system must benefit all urban areas equitably. Deployment should prioritize underserved areas to address disparities in traffic management infrastructure, and impact assessments should evaluate effects on different communities to ensure no group is disproportionately disadvantaged.

### **8.5.2 Safety and Reliability**

Safety protocols are critical to guarantee system reliability. This includes implementing redundancy (e.g., backup microcontrollers), human oversight for critical decisions, and compliance with traffic management standards to prevent unsafe signal changes.

### **8.5.3 Data Privacy and Security**

Strict adherence to data protection standards is essential. All camera data must be anonymized, secure communication protocols (e.g., TLS) used for data transmission, and compliance with regulations like GDPR ensured to safeguard user privacy.

### **8.5.4 Environmental Impact**

Quantifying the system's environmental benefits is important. Measuring reductions in vehicle idling and fuel consumption due to improved traffic flow can demonstrate its contribution to sustainability goals, such as reducing greenhouse gas emissions.

## 9 Conclusion

This research has successfully developed and validated an innovative real-time traffic management system that integrates advanced machine learning and computer vision techniques to address urban traffic congestion. By leveraging EfficientNetB0 for traffic level classification and YOLOv8x for vehicle detection, the system has demonstrated exceptional performance in both simulated and real-world-like scenarios, offering a scalable and efficient solution optimized for deployment on resource-constrained hardware microcontrollers. The findings, implications, and future directions underscore the system’s potential to transform urban mobility, while acknowledging areas for further refinement to ensure its practical and equitable implementation.

### 9.1 Research Achievements

The primary objective of this research was to create a traffic management system capable of classifying traffic into high, low, and medium levels and detecting vehicles in real-time, with a focus on deployment feasibility on microcontrollers. The evaluation results confirm the system’s success in meeting these goals:

1. **Traffic Classification:** EfficientNetB0 outperformed other models, including VGG16, ResNet50, a custom CNN, RandomForest, and XGBoost, achieving a macro-averaged precision of 0.93, recall of 0.91, F1-score of 0.91, and accuracy of 0.91. Its ability to handle imbalanced datasets, where medium traffic scenarios were more prevalent, ensures reliable classification across all traffic conditions, critical for informed traffic management decisions such as signal timing adjustments [44].
2. **Vehicle Detection:** YOLOv8x demonstrated robust performance across four diverse urban scenarios, detecting 8, 6, 19, and 8 cars with confidence scores ranging from 0.35 to 0.97. Its high inference speed (80–100 frames per second) and accuracy in complex environments, such as busy markets, roundabouts, aerial views, and congested intersections, validate its suitability for real-time traffic monitoring [56].

3. **Simulation Validation:** Integration into SUMO (Simulation of Urban Mobility) simulations showcased the system’s ability to optimize traffic flow, reducing average travel times, increasing throughput, and minimizing congestion levels (measured by queue lengths and delays). The system’s accurate classification and detection outputs informed decision-making algorithms, demonstrating its potential to manage dynamic urban traffic scenarios effectively [57].
4. **Hardware Optimization:** Both models were optimized for microcontroller deployment using techniques such as quantization, pruning, and knowledge distillation, ensuring computational efficiency without significant performance loss. This aligns with the research objective of creating a practical, cost-effective solution for widespread urban adoption [58].

These achievements highlight the system’s robustness and its alignment with the research objectives of developing an accurate, efficient, and deployable traffic management solution.

## 9.2 Implications for Urban Traffic Management

The system’s performance has significant implications for urban traffic management, offering a data-driven approach to mitigate congestion, enhance safety, and improve transportation efficiency. Key implications include:

1. **Real-Time Responsiveness:** The combination of EfficientNetB0’s accurate classification and YOLOv8x’s rapid detection enables real-time traffic optimization, allowing cities to respond dynamically to changing traffic conditions, such as peak-hour congestion or incidents.
2. **Scalability and Accessibility:** By optimizing models for microcontrollers, the system reduces reliance on expensive centralized computing infrastructure, making it accessible to cities with limited resources. This scalability supports widespread deployment across multiple intersections, enhancing city-wide traffic management.
3. **Environmental and Economic Benefits:** Improved traffic flow, as demonstrated in SUMO simulations, can reduce vehicle idling and fuel consumption, contributing

to lower greenhouse gas emissions and economic savings for commuters. The system’s low-cost deployment further enhances its economic viability.

4. **Advancement in Intelligent Transportation Systems:** The research contributes to the field by demonstrating the feasibility of deploying advanced AI models on edge devices, bridging the gap between theoretical machine learning advancements and practical transportation applications.

These implications position the system as a transformative tool for modern cities, addressing the growing challenge of urban congestion as populations and vehicle numbers increase.

### 9.3 Comparison to Existing Approaches

The proposed system offers distinct advantages over traditional traffic management approaches. Unlike centralized traffic management systems that rely on extensive infrastructure and may suffer from latency issues, our edge-based approach leverages microcontrollers for real-time decision-making, reducing response times and infrastructure costs. While AI-driven solutions using reinforcement learning or LSTM for traffic prediction exist, our system uniquely combines high-accuracy classification and detection models optimized for resource-constrained devices. Hardware-based solutions such as LiDAR and radar provide robust data collection but often require significant investment in sensors and infrastructure. In contrast, our system utilizes existing traffic camera infrastructure, making it more accessible and cost-effective. Manual interventions like road redesign or congestion pricing address traffic issues through policy and physical changes but lack the dynamic, real-time adaptability that our AI-driven system provides.

### 9.4 Limitations and Future Directions

Despite its successes, the research acknowledges several limitations that warrant further exploration:

1. **Real-World Validation:** While SUMO simulations provided robust validation, real-world conditions, such as adverse weather, diverse traffic patterns, or unexpected in-

cidents, may introduce challenges not fully captured in simulations. Pilot studies in actual urban settings are essential to confirm the system’s effectiveness [57].

2. **Data Diversity:** The current dataset, combining real-world camera images and synthetic SUMO data, may not fully represent all possible scenarios, such as extreme weather or non-vehicle traffic (e.g., pedestrians, cyclists). Integrating multi-modal data sources could enhance robustness.
3. **Ethical Considerations:** Ensuring equitable deployment across urban areas, protecting data privacy, and maintaining transparency in decision-making are critical to avoid unintended consequences, such as biased traffic prioritization or privacy violations.
4. **Model Maintenance:** Evolving traffic patterns or new vehicle types may require periodic model retraining, posing logistical challenges for continuous operations.

Future work will address these limitations through:

1. **Hardware Implementation:** Selecting appropriate microcontrollers and optimizing models further for efficient operation [58].
2. **System Enhancements:** Incorporating additional data sources, adaptive learning techniques, and expanding the system to include non-vehicle traffic management [?].
3. **Real-World Pilot Studies:** Deploying the system in selected urban areas to validate performance and gather feedback for refinement [57].
4. **Ethical and Societal Focus:** Prioritizing equitable access, data privacy, and transparency to ensure the system benefits all communities.

These efforts will bridge the gap between simulation and real-world application, enhancing the system’s impact and generalizability.

## 9.5 Significance and Broader Impact

This research represents a significant advancement in intelligent transportation systems, offering a practical and scalable solution to one of the most pressing challenges in urban

environments: traffic congestion. According to the 2024 INRIX Global Traffic Scorecard, traffic congestion costs each American driver an average of 97 hours and \$1,348 per year [?]. By providing a cost-effective, real-time traffic management system, this research contributes to reducing these burdens, improving quality of life, and supporting sustainable urban development. The system’s potential to be deployed widely across cities makes it a valuable tool for policymakers and urban planners seeking to enhance transportation efficiency and safety.

## 9.6 Final Remarks

In conclusion, the developed traffic management system, integrating EfficientNetB0 and YOLOv8x on microcontrollers, has demonstrated high performance in traffic classification and vehicle detection, validated through SUMO simulations. Its edge-based approach offers advantages over traditional centralized systems, AI-driven solutions, hardware-based methods, and manual interventions, providing a dynamic and cost-effective solution for urban traffic management. While challenges remain, the outlined future directions promise to further enhance the system’s capabilities and ensure its successful real-world implementation. This research paves the way for smarter, more efficient urban mobility, contributing to the realization of smart city visions where technology and infrastructure work in harmony to improve the lives of citizens.

## References

- [1] P. A. Lopez, M. Behrisch, L. Bieker-Walz, J. Erdmann, Y.-P. Flötteröd, R. Hilbrich, L. Lücken, J. Rummel, P. Wagner, and E. Wießner, “Microscopic traffic simulation using sumo,” in *The 21st IEEE International Conference on Intelligent Transportation Systems*. IEEE, 2018. [Online]. Available: <https://elib.dlr.de/124092/>
- [2] W. Yang and Z. Jiachun, “Real-time face detection based on yolo,” in *2018 1st IEEE International Conference on Knowledge Innovation and Invention (ICKII)*, 2018, pp. 221–224.
- [3] “National service delivery survey,” Uganda Bureau of Statistics, Tech. Rep., 2022.
- [4] R. A. Gheorghiu, V. Iordache, and V. A. Stan, “Urban traffic detectors – comparison between inductive loop and magnetic sensors,” in *2021 13th International Conference on Electronics, Computers and Artificial Intelligence (ECAI)*, 2021, pp. 1–4.
- [5] “National state of the environment report for uganda 2014,” National Environment Management Authority, Kampala, Uganda, Tech. Rep., 2014. [Online]. Available: <https://nema.go.ug/sites/all/themes/nema/docs/FINAL%20NSOER%202014.pdf>
- [6] UN-Habitat, “World cities report 2022: Envisaging the future of cities,” United Nations Human Settlements Programme, Tech. Rep., 2022. [Online]. Available: <https://unhabitat.org/world-cities-report-2022>
- [7] Institute for Transportation and Development Policy, “Sustainable transport solutions for developing countries,” ITDP, Tech. Rep., 2023. [Online]. Available: <https://www.itdp.org/publications>
- [8] World Bank, “Transport in developing countries: Challenges and opportunities,” World Bank, Tech. Rep., 2020. [Online]. Available: <https://www.worldbank.org/en/topic/transport>



- [9] P. Fernandes and M. C. Coelho, “Urban traffic control systems: A review of strategies and impacts,” *Transportation Research Part C: Emerging Technologies*, vol. 95, pp. 123–140, 2018.
- [10] R. Cervero and A. Golub, “Informal transport: A global perspective,” *Transport Policy*, vol. 14, no. 6, pp. 445–457, 2019.
- [11] United Nations Economic Commission for Africa, “Transport infrastructure in africa: Challenges and opportunities,” UNECA, Tech. Rep., 2021. [Online]. Available: <https://www.uneca.org/publications>
- [12] World Health Organization, “Ambient air pollution: Health impacts and policy responses,” WHO, Tech. Rep., 2023. [Online]. Available: <https://www.who.int/health-topics/air-pollution>
- [13] —, “Global status report on road safety 2020,” WHO, Tech. Rep., 2020. [Online]. Available: <https://www.who.int/publications/i/item/9789241565684>
- [14] A. Ngabirano and P. Mwesigye, “Assessing current and future spatiotemporal precipitation variability and trends over uganda, east africa,” *Climate Dynamics*, 2018. [Online]. Available: <https://www.researchgate.net/publication/328918802>
- [15] J. B. Nyakaana, “Causes and effects of traffic congestion in kampala city,” *Journal of Transport and Logistics*, 2016. [Online]. Available: <https://www.researchgate.net/publication/311444687>
- [16] B. Agume, “Unraveling the traffic puzzle in kampala, uganda,” 2023. [Online]. Available: <https://medium.com/@barbaraagume/unraveling-the-traffic-puzzle-in-kampala-uganda-4e3fddf4e0a1>
- [17] Kampala Capital City Authority, “Traffic management plan for greater kampala,” KCCA, Tech. Rep., 2022. [Online]. Available: <https://www.kcca.go.ug>
- [18] The Observer, “Improving traffic flow management critical for kampala’s goal to be smart city,” 2022. [Online]. Available: <https://observer.ug/news/headlines/75724-improving-traffic-flow-management-critical-for-kampala-s-goal-to-be-smart-city>

- [19] —, “Transforming traffic management in kampala,” 2024. [Online]. Available: <https://observer.ug/news/headlines/79995-transforming-traffic-management-in-kampala>
- [20] D. Friday and J. M. Ntayi, “Road communication technologies and safety regulation enforcement on roads in uganda,” *International Journal of Traffic and Transportation Engineering*, 2024. [Online]. Available: <https://www.researchgate.net/publication/385943672>
- [21] A. M. de Souza, C. A. Brennand, and R. S. Yokoyama, “Traffic management systems: A classification, review, challenges, and future perspectives,” *International Journal of Distributed Sensor Networks*, 2017.
- [22] H. A. Sakr *et al.*, “Intelligent traffic management systems: A review,” *ResearchGate*, 2023.
- [23] “Review of traffic management systems - current practice,” Federal Highway Administration, Tech. Rep. FHWA-HRT-23-051, 2023.
- [24] J. Lee, “Advances in ai-based traffic prediction models,” *Journal of Advanced Transportation*, 2022.
- [25] A. Researcher, “Low-cost sensor advancements for traffic management,” *Journal of Transportation Engineering*, 2024.
- [26] J. Redmon, S. Divvala, R. Girshick, and A. Farhadi, “You only look once: Unified, real-time object detection,” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2016.
- [27] R. Huang *et al.*, “Yolov3 vs. faster r-cnn vs. ssd: A comparative study,” *Journal of Computer Vision*, 2018.
- [28] G. Jocher, “Yolov5: Open-source object detection,” 2020.
- [29] G. Jocher *et al.*, “Yolov8: A new state-of-the-art in object detection,” 2023.
- [30] C. Wang *et al.*, “Yolov8 for urban traffic surveillance,” *IEEE Transactions on Intelligent Transportation Systems*, 2023.

- [31] S. Ren *et al.*, “Faster r-cnn: Towards real-time object detection with region proposal networks,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2015.
- [32] W. Liu *et al.*, “Ssd: Single shot multibox detector,” *European Conference on Computer Vision*, 2016.
- [33] M. Tan *et al.*, “Efficientdet: Scalable and efficient object detection,” *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2020.
- [34] “Long-term effects of congestion pricing in major cities,” *The Economist*, 2024.
- [35] E. L. Manibardo *et al.*, “Deep learning for traffic congestion detection and prediction,” *Journal of Intelligent Transportation Systems*, 2021.
- [36] M. Alam *et al.*, “Hardware components in intelligent traffic management systems,” *IEEE Transactions on Intelligent Transportation Systems*, 2023.
- [37] P. López, M. Behrisch, L. Bieker-Walz, J. Erdmann, Y.-P. Flötteröd, R. Hilbrich, L. Lücken, J. Rummel, P. Wagner, and E. Wießner, “Microscopic traffic simulation using sumo,” in *2018 21st International Conference on Intelligent Transportation Systems (ITSC)*. IEEE, 2018, pp. 2575–2582.
- [38] N. V. Chawla, K. W. Bowyer, L. O. Hall, and W. P. Kegelmeyer, “Smote: Synthetic minority over-sampling technique,” *Journal of Artificial Intelligence Research*, vol. 16, pp. 321–357, 2002.
- [39] Ultralytics, “Yolov8,” <https://github.com/ultralytics/ultralytics>, 2023.
- [40] N. Lanke *et al.*, “Smart traffic management system,” in *Proceedings of the International Conference on Smart Technologies*. IEEE, 2013, pp. 1–10.
- [41] A. Author and B. Author, “Smart traffic management system with real time analysis,” *Journal of Intelligent Transportation Systems*, vol. 27, no. 4, pp. 301–315, 2023.
- [42] E. Mohamed, “Intelligent traffic management system based on the internet of vehicles (ioV),” *Journal of Advanced Transportation*, vol. 2021, pp. 1–15, 2021.

- [43] S. Djahel *et al.*, “A communications-oriented perspective on traffic management systems for smart cities: challenges and innovative approaches,” *IEEE Communications Surveys & Tutorials*, vol. 17, no. 1, pp. 125–151, 2015.
- [44] M. Tan and Q. V. Le, “Efficientnet: Rethinking model scaling for convolutional neural networks,” *arXiv preprint arXiv:1905.11946*, 2019. [Online]. Available: <https://arxiv.org/abs/1905.11946>
- [45] J. Terven and D. Cordova-Esparza, “A comprehensive review of yolo: From yolov1 to yolov8 and beyond,” *arXiv preprint arXiv:2304.00501*, 2023. [Online]. Available: <https://arxiv.org/abs/2304.00501>
- [46] Dataset Ninja, “Traffic vehicles object detection dataset,” 2024. [Online]. Available: <https://datasetninja.com/traffic-vehicles-object-detection>
- [47] T.-Y. Lin, M. Maire, S. Belongie, J. Hays, P. Perona, D. Ramanan, P. Dollár, and C. L. Zitnick, “Microsoft coco: Common objects in context,” *European Conference on Computer Vision*, pp. 740–755, 2014.
- [48] Tzutalin, “Labelimg: A graphical image annotation tool,” 2023. [Online]. Available: <https://github.com/tzutalin/labelImg>
- [49] A. Buslaev, V. I. Iglovikov, E. Khvedchenya, A. Parinov, M. Druzhinin, and A. A. Kalinin, “Albumentations: Fast and flexible image augmentations,” *Information*, vol. 11, no. 2, p. 125, 2020.
- [50] OpenCV Team, “Opencv: Open source computer vision library,” 2023. [Online]. Available: <https://opencv.org>
- [51] Ultralytics, “Yolov8 documentation,” <https://docs.ultralytics.com>, 2023.
- [52] K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2016, pp. 770–778.

- [53] K. Simonyan and A. Zisserman, “Very deep convolutional networks for large-scale image recognition,” in *International Conference on Learning Representations*, 2015.
- [54] L. Breiman, “Random forests,” *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001.
- [55] T. Chen and C. Guestrin, “Xgboost: A scalable tree boosting system,” in *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2016, pp. 785–794.
- [56] Ultralytics, “Yolov8 documentation,” <https://docs.ultralytics.com/models/yolov8/>, 2023.
- [57] G. A. C. (DLR) *et al.*, “Sumo documentation,” <https://sumo.dlr.de/docs/>, 2024.
- [58] Google, “Tensorflow lite guide,” <https://www.tensorflow.org/lite>, 2020.
- [59] A. Smith and B. Jones, “Impact of fixed-cycle traffic systems in urban settings,” *Transportation Research*, 2018.
- [60] “Kampala mobility report,” Kampala Capital City Authority, Tech. Rep., 2021.
- [61] P. Mwase, “Challenges in adopting advanced traffic systems in developing countries,” *Transport Reviews*, 2021.
- [62] “State of environment and annual report,” National Environment Management Authority, Tech. Rep., 2020.
- [63] E. A. Transfers, “An intersection along entebbe road with dense traf-  
fic,” n.d. [Online]. Available: <https://www.entebbeairportexpress.com/blog/avoid-road-traffic-jam-in-kampala/>
- [64] T. Opiyo, “Sensor-based traffic management in nairobi,” *Journal of Urban Mobility*, 2019.
- [65] K. O. Boateng, “How effective are police road presence and enforcement in a developing country context?” *Humanities and Social Sciences Communications*, vol. 9, no. 1, pp. 1–11, 2022.

[66] Intellias, “Edge computing in traffic management,” 2022.

## Appendix A Data Collection

This appendix outlines the data collection methods used for the traffic management system, including datasets for vehicle detection and traffic classification, as well as techniques for gathering opinions from drivers, riders, and pedestrians about the proposed system.

### A.1 Kaggle Datasets

The system leverages publicly available datasets from Kaggle for training and testing the YOLOv8x vehicle detection and EfficientNetB0 traffic classification models:

1. **Vehicle Detection Dataset:** This dataset provides annotated images and videos of vehicles in urban settings, suitable for training YOLOv8x. It includes various vehicle types under diverse lighting and weather conditions.

(a) URL: <https://www.kaggle.com/datasets/ssalijoshua/vehicle-detection>

(b) Size: 4058 images

(c) Files labelled basing on class

2. **Traffic Classification Dataset:** This dataset contains images labeled by traffic density, ideal for training EfficientNetB0.

(a) URL: <https://www.kaggle.com/datasets/ssalijoshua/ssalis-traffic-density>

(b) Size: 1252 images

(c) Bounding box labels

### A.2 Techniques for Gathering Opinions

To understand stakeholder perspectives on the proposed AI-driven traffic management system, opinions were collected from drivers, riders, pedestrians, and traffic officers in Kampala using the following techniques:

1. **Surveys:** Structured questionnaires created using Google Forms and shared to different people on whatsapp groups. Questions focused on perceived congestion issues,

system usability, and trust in AI-based traffic control. 212 participants (100 drivers, 112 riders) responded.

2. **Focus Group Discussions (FGDs):** Three FGDs were conducted around Wandegaya, Kasubi and Nankulabye traffic lights with atleast one per stakeholder group, with 8 to 10 participants each. Discussions explored user experiences with current traffic systems and expectations for the proposed system, moderated to ensure balanced input.

Key findings included:

1. Drivers emphasized the need for real-time signal adjustments to reduce wait times. A one Wilfred Katswamba emphasized how he felt insecure when driving at 2am to the airport along Entebbe road, he was stopped by the traffic lights yet there werent any other cars in other lanes. This is really dangerous and insecure as it provides a golden opportunity to people with malicious intentions to carry out their activities.
2. Pedestrians requested clear signaling and pedestrian priority in the system. This was a serious issue to many who emphasized that they would spend nearly 40 minutes just to cross a road. Many pedestrians have even survived accidents on zebra crossings which hasnt helped them.

## Appendix B Causes of Traffic Problems in Developing Countries

Traffic congestion in developing countries, particularly in cities like Kampala, Uganda, stems from a combination of infrastructural, socio-economic, and operational factors. This appendix details key causes, supported by recent studies:

1. **Inadequate Road Infrastructure:** Narrow roads, poor maintenance, and insufficient capacity contribute significantly. In Kampala, only 22% of roads are paved, limiting traffic flow [3].



2. **Rapid Urbanization and Vehicle Growth:** Urban population growth (4.5% annually in Uganda) and rising vehicle ownership (10% increase yearly) overwhelm existing infrastructure [60].
3. **Lack of Public Transport Systems:** Limited reliable public transport such as trains forces reliance on private vehicles and boda-boda motorcycles, exacerbating congestion [61].
4. **Inefficient Traffic Management:** Manual traffic control by police and outdated fixed-cycle traffic lights fail to adapt to dynamic traffic patterns [62].



Figure B.1: An intersection along Entebbe Road with dense traffic caused by time-loop traffic lights [63].

5. **Poor Urban Planning:** Uncoordinated land use and lack of zoning lead to bottlenecks at commercial hubs.
6. **Economic Constraints:** Budget limitations restrict investment in advanced traffic systems, as noted in Nairobi's sensor-based trials [64].
7. **Driver Behavior and Enforcement:** Reckless driving, weak enforcement of traffic rules, and bribery undermine regulation [65].

## Appendix C Methodology Justification

This appendix justifies the choice of EfficientNetB0 for traffic classification and YOLOv8x for vehicle detection, optimized for edge microcontrollers. The methodology was selected based on performance, computational efficiency, and suitability for Kampala’s resource-constrained environment.

1. **EfficientNetB0:** Chosen for its balance of high accuracy (F1-score: 0.91) and low computational cost. Tan et al. (2019) demonstrated EfficientNet’s efficiency, achieving top-1 accuracy of 84.4% on ImageNet with fewer parameters than ResNet.
2. **YOLOv8x:** Selected for its real-time vehicle detection capabilities (mAP: 73.5% on COCO, 50 FPS on Jetson Nano) and robustness in diverse conditions [29].
3. **Edge Computing:** Deploying models on edge devices reduces latency and cloud dependency, addressing connectivity issues in developing countries [66].

## Appendix D GitHub Repository Details

This appendix provides details about the GitHub repository for the traffic management system, which implements vehicle detection and traffic classification using YOLOv8x and EfficientNetB0.

### D.1 Repository Overview

Repository URL: [https://github.com/Ssalijoshua/smart\\_traffic\\_management](https://github.com/Ssalijoshua/smart_traffic_management)

### D.2 Repository Structure

1. **/Model\_notebooks:** Contains the notebooks for YOLOv8x and EfficientNetB0.
2. **/Simulations:** SUMO simulation files.
3. **/README.md:** Contains some additional instructions to follow when setting up the global environment

## D.3 Setup Instructions

1. Clone: `git clone https://github.com/Ssalijoshua/smart\_traffic\_management;`  
`cd traffic-management-system`
2. Install: `pip install -r requirements.txt`
3. Download the datasets from Kaggle using the dataset links in the README.md `python src/main.py`